0   Instructions (Total 50pts)

Homework is due Monday, March 18, 2024 at 23:59pm Central Time. Please refer to `https://courses.grainger.illinois.edu/cs446/sp2024/homework/hw/index.html` for course policy on homeworks and submission instructions.

1   Neural networks for simple functions (10pts)

In this problem you will be hand designing simple neural networks to model specific functions. In each case assume $x \in \mathbb{R}$ and provide appropriate weights $w_0, w_1, b_0 \in \mathbb{R}^d$ for some d. In other words, each neural network has one input neuron, d hidden neurons, and one output neuron.

1. (1pts) $w_0, w_1 \in \mathbb{R}^2$ such that $f(x) = w_1^T \sigma(w_0 x) = x, \forall x \in \mathbb{R}$ where $\sigma = \text{ReLU}$ (note: $w_0 x$ here is a vector-scalar product: $\mathbb{R}^2 \times \mathbb{R} \to \mathbb{R}^2$, e.g. $[0,1]^T x = [0,x]^T$)

2. (1pts) $w_0, w_1, b_0 \in \mathbb{R}^2$ such that $f(x) = w_1^T \sigma(w_0 x + b_0) = mx + b, \forall x \in \mathbb{R}$ where $\sigma = ReLU$

3. (1pts) $w_0, w_1 \in \mathbb{R}$ such that $f(x) = w_1^T \sigma(w_0 x) = b, \forall x \in \mathbb{R}$ where $\sigma = sigmoid$

4. (2pts) $w_0, w_1, b_0 \in \mathbb{R}^3, \sigma = ReLU$, and $f(x) = w_1^T \sigma(w_0 x + b_0) = \begin{cases} 0 & x \leq -2 \\ 3x + 6 & -2 \leq x \leq 0 \\ -3x + 6 & 0 \leq x \leq 2 \\ 0 & x \geq 2 \end{cases}$

5. (1pts) Do there exist $w_0, w_1 \in \mathbb{R}^d$ such that $f(x) = w_1^T \sigma(w_0 x) = x^2, \forall x \in \mathbb{R}$ where $\sigma = ReLU$. Explain why or why not.

6. (2pts) For any $\epsilon > 0$ do there exist fixed finite dimensional $w_0, w_1, b_0 \in \mathbb{R}^d$ such that $|f(x) - x^2| = |w_1^T \sigma(w_0 x + b_0) - x^2| < \epsilon, \forall x$ where $\sigma = ReLU$? Explain why or why not. If yes describe a procedure to construct w_0, w_1, b_0 at a high level.

7. (2pts) For any $\epsilon > 0$ do there exist fixed finite dimensional w_0, w_1, b_0 such that $|f(x) - x^2| = |w_1^T \sigma(w_0 x + b_0) - x^2| < \epsilon$ where $\sigma = ReLU$ and $x \in [0,1]$? Explain why or why not. If yes describe a procedure to construct w_0, w_1, b_0 at a high level.

2 Backpropagation through time (BPTT) (8pts)

1. (2pts) Consider recurrent model with recurrence $h_t = w(x_t + h_{t-1})$, and $h_0 = 0$. Draw the computation graph for unrolling this RNN for 3 timesteps. You should include 3 inputs x_1, x_2, x_3, the outputs h_0, h_1, h_2, h_3, and refer to intermediate values $y_i = (x_i + h_{i-1})$.

2. (1pts) Compute the forward pass in the above model, i.e., compute $h_0, h_1, h_2, h_3, y_1, y_2, y_3$ as functions of w, x_1, x_2, x_3.

3. (2pts) Compute $\frac{\partial h_3}{\partial w}$ in the above model using a backwards pass and your values from the forward pass. HINT: you can directly take the derivative of h_3 from the previous question to check your work, but you must show work for a proper backprop including intermediate partial derivatives.

4. (2pts) For sequence $z = [x_1, ..., x_T]$ Consider recurrent model $f(z) = h_T$, where $h_t = w\sigma(x_t + h_{t-1})$, and $h_0 = 0$. Derive $\frac{\partial f}{\partial h_1} = \frac{\partial h_T}{\partial h_1}$. You may use σ' to represent the derivative of nonlinearity σ.

5. (1pts) Using your previous answer as justification, explain the exploding/vanishing gradient problem in RNNs.

3   Transformers (7pts)

In self attention we compute the output vector $c = \sum_{i=1}^{n} \alpha_i v_i$ for n value vectors $v_i \in \mathbb{R}^d$ and attention weights $\alpha_i = \frac{exp(k_i^T q)}{\sum_j exp(k_j^T q)}$ computed as a softmax over query vector $q \in \mathbb{R}^d$ and n keys $k_i \in \mathbb{R}^d$.

1. (1pts) Can α_i be infinitely large? Can it be zero? Explain your answer.

2. (1pts) Describe q and k_i such that $\alpha_i \gg \alpha_j$ for some i and all $j \neq i$. What is c in this case? Explain intuitively what such a case represents in the context of attention.

3. (1pts) Assuming all key vectors are orthonormal, $k_i^T k_j = 0$ and $||k_i|| = 1$, give a q such that $c \approx \frac{1}{2}(v_1 + v_2)$.

4. (2pts) The dot product is one way of measuring similarity between the query and the keys. We can replace it with any similarity, e.g., a kernel. Show that you can rewrite the dot product softmax similarity, $\exp(q^T k_i)$, with the RBF kernel, $\exp(\frac{-||q-k_i||_2^2}{2\sigma})$. You may assume the keys all have norm 1, $||k_i||_2 = 1$.

5. (2pts) In general we have one query per element in the sequence of length n. We can rewrite such an operation as $attention(Q, K, V) = softmax(QK^T)V$ for 3 (n, d) matrices. How does computing the output scale with n, the sequence length? Let $P = softmax(QK^T)$ and assume P has rank k and we know its SVD. Show that attention can be computed in $\mathcal{O}(nkd)$ time.

4 Resnet (12pts)

In this problem, you will implement a simplified ResNet. You do not need to change arguments which are not mentioned here (but you of course could try and see what happens).

1. (5pts) Implement a class `Block`, which is a building block of ResNet.

   The input to `Block` is of shape (N, C, H, W), where N denotes the batch size, C denotes the number of channels, and H and W are the height and width of each channel. For each data example $\boldsymbol{x}$ with shape (C, H, W), the output of `block` is

   $$\texttt{Block}(\boldsymbol{x}) = \sigma_r(\boldsymbol{x} + f(\boldsymbol{x})),$$

   where σ_r denotes the ReLU activation, and $f(\boldsymbol{x})$ also has shape (C, H, W) and thus can be added to $\boldsymbol{x}$. In detail, f contains the following layers.

   (a) A `Conv2d` with C input channels, C output channels, kernel size 3, stride 1, padding 1, and no bias term.

   (b) A `BatchNorm2d` with C features.

   (c) A ReLU layer.

   (d) Another `Conv2d` with the same arguments as (a) above.

   (e) Another `BatchNorm2d` with C features.

   Because 3×3 kernels and padding 1 are used, the convolutional layers do not change the shape of each channel. Moreover, the number of channels are also kept unchanged. Therefore $f(\boldsymbol{x})$ does have the same shape as $\boldsymbol{x}$.

   Additional instructions are given in doscstrings in `hw3.py`.

   Suggested Library routines: `torch.nn.Conv2d` and `torch.nn.BatchNorm2d`.

   Remark: Use `bias=False` for the `Conv2d` layers.

2. (2pts) Explain why a `Conv2d` layer does not need a bias term if it is followed by a `BatchNorm2d` layer.

3. (5pts) Implement a (shallow) `ResNet` consists of the following parts:

   (a) A `Conv2d` with 1 input channel, C output channels, kernel size 3, stride 2, padding 1, and no bias term.

   (b) A `BatchNorm2d` with C features.

(c) A ReLU layer.

(d) A `MaxPool2d` with kernel size 2.

(e) A `Block` with C channels.

(f) An `AdaptiveAvgPool2d` which for each channel takes the average of all elements.

(g) A `Linear` with C inputs and 10 outputs.

Additional instructions are given in doscstrings in `hw3.py`.

Suggested Library routines: `torch.nn.Conv2d`, `torch.nn.BatchNorm2d`, `torch.nn.MaxPool2D`, `torch.nn.AdaptiveAvgPool2d` and `torch.nn.Linear`.

Remark: Use `bias=False` for the `Conv2d` layer.

Remark: It is not required for you to train the model for the purpose of this problem.

5 Image overfitting (13pts)

In this problem, you will be overfitting a neural network to an image. In other words, for a grayscale image M you will train a continuous function $f_\theta(x, y) : \mathbb{R}^2 \to \mathbb{R}$ such that $f(x, y)$ approximates $M_{x,y}$ for all image coordinate (x, y). Your dataset will consist of all coordinate-pixel pairs for a *single image*, and your task will be to experiment with different activation functions. You can run the code in this problem with
`python siren.py -b batch_size -e total_epochs -a [sin, tanh, relu]`

1. Finish the implementation of the MyDataset class - namely implement the __get_item__(self, idx) function which returns the data element corresponding to the given index.

2. Implement the SingleLayer class (used by class Siren), a simple Linear layer followed by an activation. You should initialize the weights of the first layer to be uniform $[\frac{-1}{input\_features}, \frac{1}{input\_features}]$ and every other layer to be uniform$[-\frac{\sqrt{6/hidden\_features}}{\omega}, \frac{\sqrt{6/hidden\_features}}{\omega}]$. $\omega = 30$ for sin activation and $\omega = 1$ for other activations. In the forward(self, input) function you should multiply the layer output by ω before applying the activation.

3. Set your learning rate appropriately in train()

4. Implement the main loop of the train() function, which trains the Siren model with SGD over batches from the dataloader.

5. (9pts) Run experiments with ReLU, Tanh, and Sin activations. For each activation report one successful experiment including:

   (a) total number of epochs

   (b) learning rate

   (c) batch size

   (d) plot a loss curve

   (e) plot model output images, gradients, and laplacians (generated by model_results())

6. (1pts) Provide a mathematical explanation for why the Laplacian looks as it does for ReLU activation.

7. (1pts) When your loss plateaus, what is something you can do to the learning rate to keep improving your model? What is this strategy called?

8. (2pts) Describe what happens if you do not manually set the initial weights and what happens when you don't normalize the input coordinates? You should report the results for an experiment in each of these settings.