

0 Instructions

Homework is due Tuesday, April 2, 2024 at 23:59pm Central Time. Please refer to <https://courses.grainger.illinois.edu/cs446/sp2024/homework/hw/index.html> for course policy on homeworks and submission instructions.

Reminder: Answers must be typeset. L^AT_EX and other methods of typesetting math are accepted.

1 PCA: 6pts

- (1pts) Recall that PCA finds a direction w in which the projected data has highest variance by solving the following program:

$$\max_{w: \|w\|^2=1} w^T \Sigma w. \quad (1)$$

Here, Σ is a covariance matrix. You are given a dataset of two 2-dimensional points (1, 3) and (4, 7). Draw the two data points on the 2D plane. What is the first principal component w of this dataset?

- (3pts) Now you are given a dataset of four points (2, 0), (2, 2), (6, 0) and (6, 2). Given this dataset, derive the covariance matrix Σ in Eq.1. Then plot the centralized data with the first and the second principal components in one figure. **Include the plot in your written submission.**
- (2pts) What is the optimal w and the optimal value of the program in Eq.1 given

$$\Sigma = \begin{bmatrix} 12 & 0 & 0 & 0 \\ 0 & 6 & 0 & 0 \\ 0 & 0 & 20 & 0 \\ 0 & 0 & 0 & 10 \end{bmatrix}.$$

Give your justification.

2 Basics in Information Theory: 7pts

Let X be a discrete variable, and P, Q be two probability distributions over X . Define a new random variable X' as follows:

$$X' = \begin{cases} X \sim P & \text{if } B = 1, \\ X \sim Q & \text{if } B = 0, \end{cases}$$

where $B \in \{0, 1\}$ is an independent and Bernoulli distribution over $\{0, 1\}$ with the parameter λ , such that $\Pr(B = 1) = \lambda = 1 - \Pr(B = 0)$.

1. (2pts) Derive and represent the mixture distribution $\Pr(X' = x)$ in terms of $P(x)$ and $Q(x)$.
2. (5pts) Show that $I(X'; B) = D_\lambda(P\|Q)$, where $D_\lambda(P\|Q)$ is the λ -divergence between P and Q , i.e., $D_\lambda(P\|Q) = \lambda D_{\text{KL}}(P\|\lambda P + (1 - \lambda)Q) + (1 - \lambda) D_{\text{KL}}(Q\|\lambda P + (1 - \lambda)Q)$. Note that by setting $\lambda = 0.5$, the λ -divergence gives the Jensen-Shannon divergence.

3 k-Means with Soft Assignments: 10pts

Consider the following exemplar-based, hard-assignment form as the objective of k-Means for K clusters and n data points $x^{(i)}$ for $i = 1, \dots, n$:

$$\min_{\mu_1, \dots, \mu_K} \sum_{i=1}^n \min_k \|x^{(i)} - \mu_k\|_2^2 = \min_{\mu_1, \dots, \mu_K} \min_{\substack{A \in \{0,1\}^{n \times K} \\ A \cdot \mathbf{1}_K = \mathbf{1}_n}} \sum_{i=1}^n \sum_{k=1}^K A_{ik} \|x^{(i)} - \mu_k\|_2^2, \quad (2)$$

where μ_k denotes the center for the k -th cluster, the matrix $A \in \{0, 1\}^{n \times K}$ indicates the hard assignment of each data point to the clusters, and $A \cdot \mathbf{1}_K = \mathbf{1}_n$, which tells us that each row of A has one 1 with all remaining elements as 0, i.e., $\sum_{k=1}^K A_{ik} = 1, \forall i$.

We extend this setting to soft assignment by designing the matrix $A \in [0, 1]^{n \times K}$, and the objective becomes:

$$\min_{\mu_1, \dots, \mu_K} \min_{\substack{A \in [0,1]^{n \times K} \\ A \cdot \mathbf{1}_K = \mathbf{1}_n}} \sum_{i=1}^n \sum_{k=1}^K A_{ik} \|x^{(i)} - \mu_k\|_2^2. \quad (3)$$

1. (3pts) Show that the following holds:

$$\min_{\mu_1, \dots, \mu_K} \min_{\substack{A \in [0,1]^{n \times K} \\ A \cdot \mathbf{1}_K = \mathbf{1}_n}} \sum_{i=1}^n \sum_{k=1}^K A_{ik} \|x^{(i)} - \mu_k\|_2^2 \leq \min_{\mu_1, \dots, \mu_K} \min_{\substack{A \in \{0,1\}^{n \times K} \\ A \cdot \mathbf{1}_K = \mathbf{1}_n}} \sum_{i=1}^n \sum_{k=1}^K A_{ik} \|x^{(i)} - \mu_k\|_2^2$$

Hint: Note that $\{0, 1\}^{n \times K}$ can be seen as a subset of $[0, 1]^{n \times K}$.

2. (5pts) Show that the following also holds:

$$\min_{\mu_1, \dots, \mu_K} \min_{\substack{A \in \{0,1\}^{n \times K} \\ A \cdot \mathbf{1}_K = \mathbf{1}_n}} \sum_{i=1}^n \sum_{k=1}^K A_{ik} \|x^{(i)} - \mu_k\|_2^2 \geq \min_{\mu_1, \dots, \mu_K} \min_{\substack{A \in \{0,1\}^{n \times K} \\ A \cdot \mathbf{1}_K = \mathbf{1}_n}} \sum_{i=1}^n \sum_{k=1}^K A_{ik} \|x^{(i)} - \mu_k\|_2^2$$

Hint: You may use the fact that $\|x^{(i)} - \mu_k\|_2^2 \geq \min_l \|x^{(i)} - \mu_l\|_2^2$, for any i and k .

3. (2pts) Show that the soft assignment problem introduced in this problem (Eq. 3) corresponds to a globally optimal hard assignment.

4 Bernoulli Mixture Model: 18pts

Extended from the Gaussian mixture model introduced in the lecture, we explore the Bernoulli mixture model in this problem. We represent the dataset as $X \in \{0, 1\}^{n \times d}$ and each data instance is a set of d independent binary random variables $x^{(i)} = \{x_1^{(i)}, x_2^{(i)}, \dots, x_d^{(i)}\}$ and the probability that $x^{(i)}$ is generated from the k -th Bernoulli distributions is calculated as:

$$\Pr(x^{(i)} | \mu_k) = \prod_{j=1}^d \mu_k^{x_j^{(i)}} (1 - \mu_k)^{(1-x_j^{(i)})},$$

where μ_k is the mean of the k -th Bernoulli distribution.

We consider K mixed Bernoulli distributions and introduce the auxiliary/latent variable $z_{ik} \in \{0, 1\}$ with $\sum_{k=1}^K z_{ik} = 1 \forall i$ as the assignment for $x^{(i)}$ to the k -th Bernoulli distribution. Also, we have $\Pr(z_{ik} = 1) = \pi_k$ and $\Pr(x^{(i)} | z_{ik} = 1) = \Pr(x^{(i)} | \mu_k)$.

1. (5pts) Derive the log-likelihood $\log \Pr(x^{(i)}, z_i | \pi, \mu)$.
2. (5pts) In the **expectation** step, derive the update step for the assignment (posterior) $z_{ik}^{\text{new}} = \Pr(z_{ik} = 1 | x^{(i)})$.
3. (8pts) In the **maximization** step, derive the update step for the model parameter, i.e., μ_k^{new} and π_k^{new} .

5 Variational Autoencoder (VAE): 19pts

In this problem you will implement a Variational Autoencoder (VAE) to model points sampled from an unknown distribution. This will be done by constructing an encoder network and a decoder network. The encoder network $f_{\text{enc}} : X \subset \mathbb{R}^2 \rightarrow \mathbb{R}^h \times \mathbb{R}^h$ takes as input a point \mathbf{x} from the input space and outputs parameters (μ, ξ) where $\xi = \log \sigma^2$. The decoder

network $f_{\text{dec}} : \mathbb{R}^h \rightarrow \mathbb{R}^2$ takes as input a latent vector $\mathbf{z} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\sigma}^2)$ and outputs an element $\hat{\mathbf{x}} \in \mathbb{R}^2$ that we would hope is similar to members of the input space X . You will train this model by minimizing the (regularized) empirical risk

$$\hat{\mathcal{R}}_{\text{VAE}}(f) = \frac{1}{n} \sum_{i=1}^n \ell(\hat{\mathbf{x}}_i, \mathbf{x}_i) + \lambda \text{KL}(\mathcal{N}(\boldsymbol{\mu}(\mathbf{x}_i), \exp(\boldsymbol{\xi}(\mathbf{x}_i)/2)), \mathcal{N}(0, I)).$$

Particularly, we have

$$\text{KL}(\mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma}), \mathcal{N}(0, I)) = -\frac{1}{2} \left[h + \sum_{j=1}^h (\log \sigma_j^2 - \mu_j^2 - \sigma_j^2) \right],$$

1. (3pts) Use the empirical risk discussed above to implement a VAE in the class `VAE`. Use ReLU activations between each layer, except on the last layer of the decoder use sigmoid. Use the Adam optimizer to optimize in the `step()` function. Make use of the PyTorch library for this. Use `torch.optim.Adam()`, there is no need to implement it yourself. Please refer to the docstrings in `hw4.py` for more implementation details.
2. (5pts) Implement the `fit` function using the `step()` function from the `VAE` class. See the docstrings in `hw4.py` for more details.
3. (11pts) Fit a VAE on the data generated by `generate_data` in `hw4.utils.py`. Use a learning rate $\eta = 0.01$, latent space dimension $h = 6$, KL-divergence scaling factor $\lambda = 5 \times 10^{-5}$, and train for 8000 iterations. Use least squares as the loss, that is, let $\ell(\mathbf{x}, \hat{\mathbf{x}}) = \|\mathbf{x} - \hat{\mathbf{x}}\|_2^2$. Include separate plots of each of the following with a legend in your written submission:

- (a) Your empirical risk $\hat{\mathcal{R}}_{\text{VAE}}$ on the training data vs iteration count;
- (b) The data points $(\mathbf{x})_{i=1}^n$ along with their encoded and decoded approximations $\hat{\mathbf{x}}$;
- (c) The data points $(\mathbf{x})_{i=1}^n$ along with their encoded and decoded approximations $\hat{\mathbf{x}}$, and n generated points $f_{\text{dec}}(\mathbf{z})$ where $\mathbf{z} \sim \mathcal{N}(0, I)$.

After you are done training, save your neural network to a checkpoint file using `torch.save(model.cpu().state_dict(), "vae.pb")`. You will submit this checkpoint file "vae.pb" to the autograder with your code submission.