

# CS 446/ECE 449: Machine Learning

---

## Lecture 22: Diffusion Models

Han Zhao  
04/09/2024



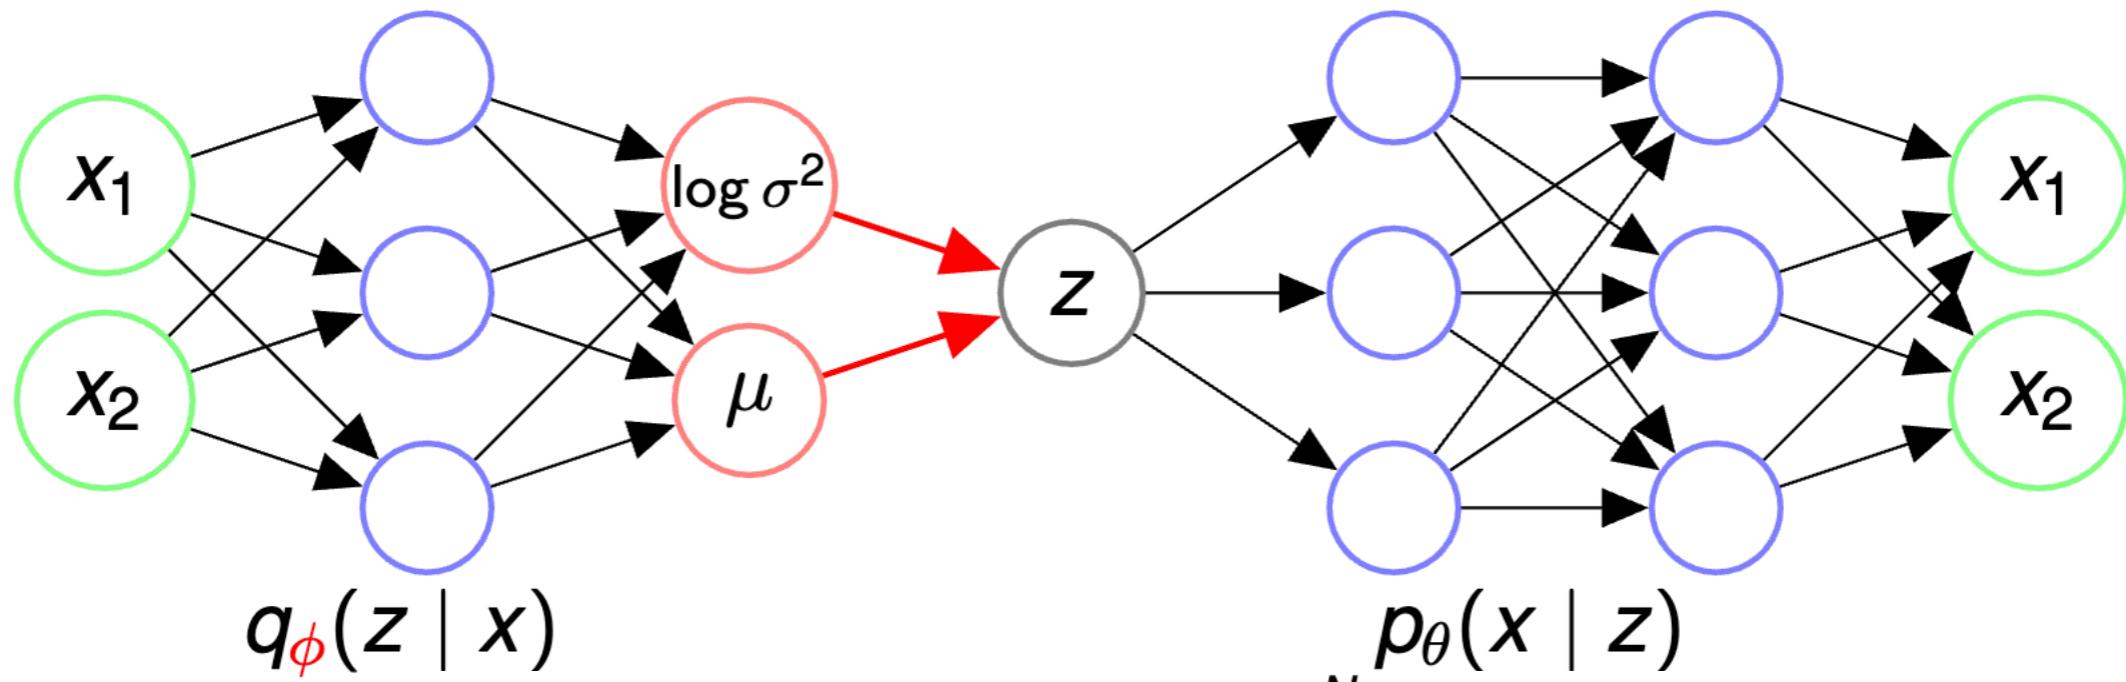
# Reference for Diffusion Models

---

- “Deep Learning: Foundations and Concepts”, Chapter 20
- “Tutorial on Diffusion Models for Imaging and Vision”, Stanley H. Chan, <https://arxiv.org/abs/2403.18103>
- “Denoising Diffusion Probabilistic Models”, <https://arxiv.org/abs/2006.11239>
- “Variational Diffusion Models”, <https://arxiv.org/abs/2107.00630>

# Recap: Variational Auto-encoder

Recall from Lecture 20 we introduced the variational auto-encoder (VAE):



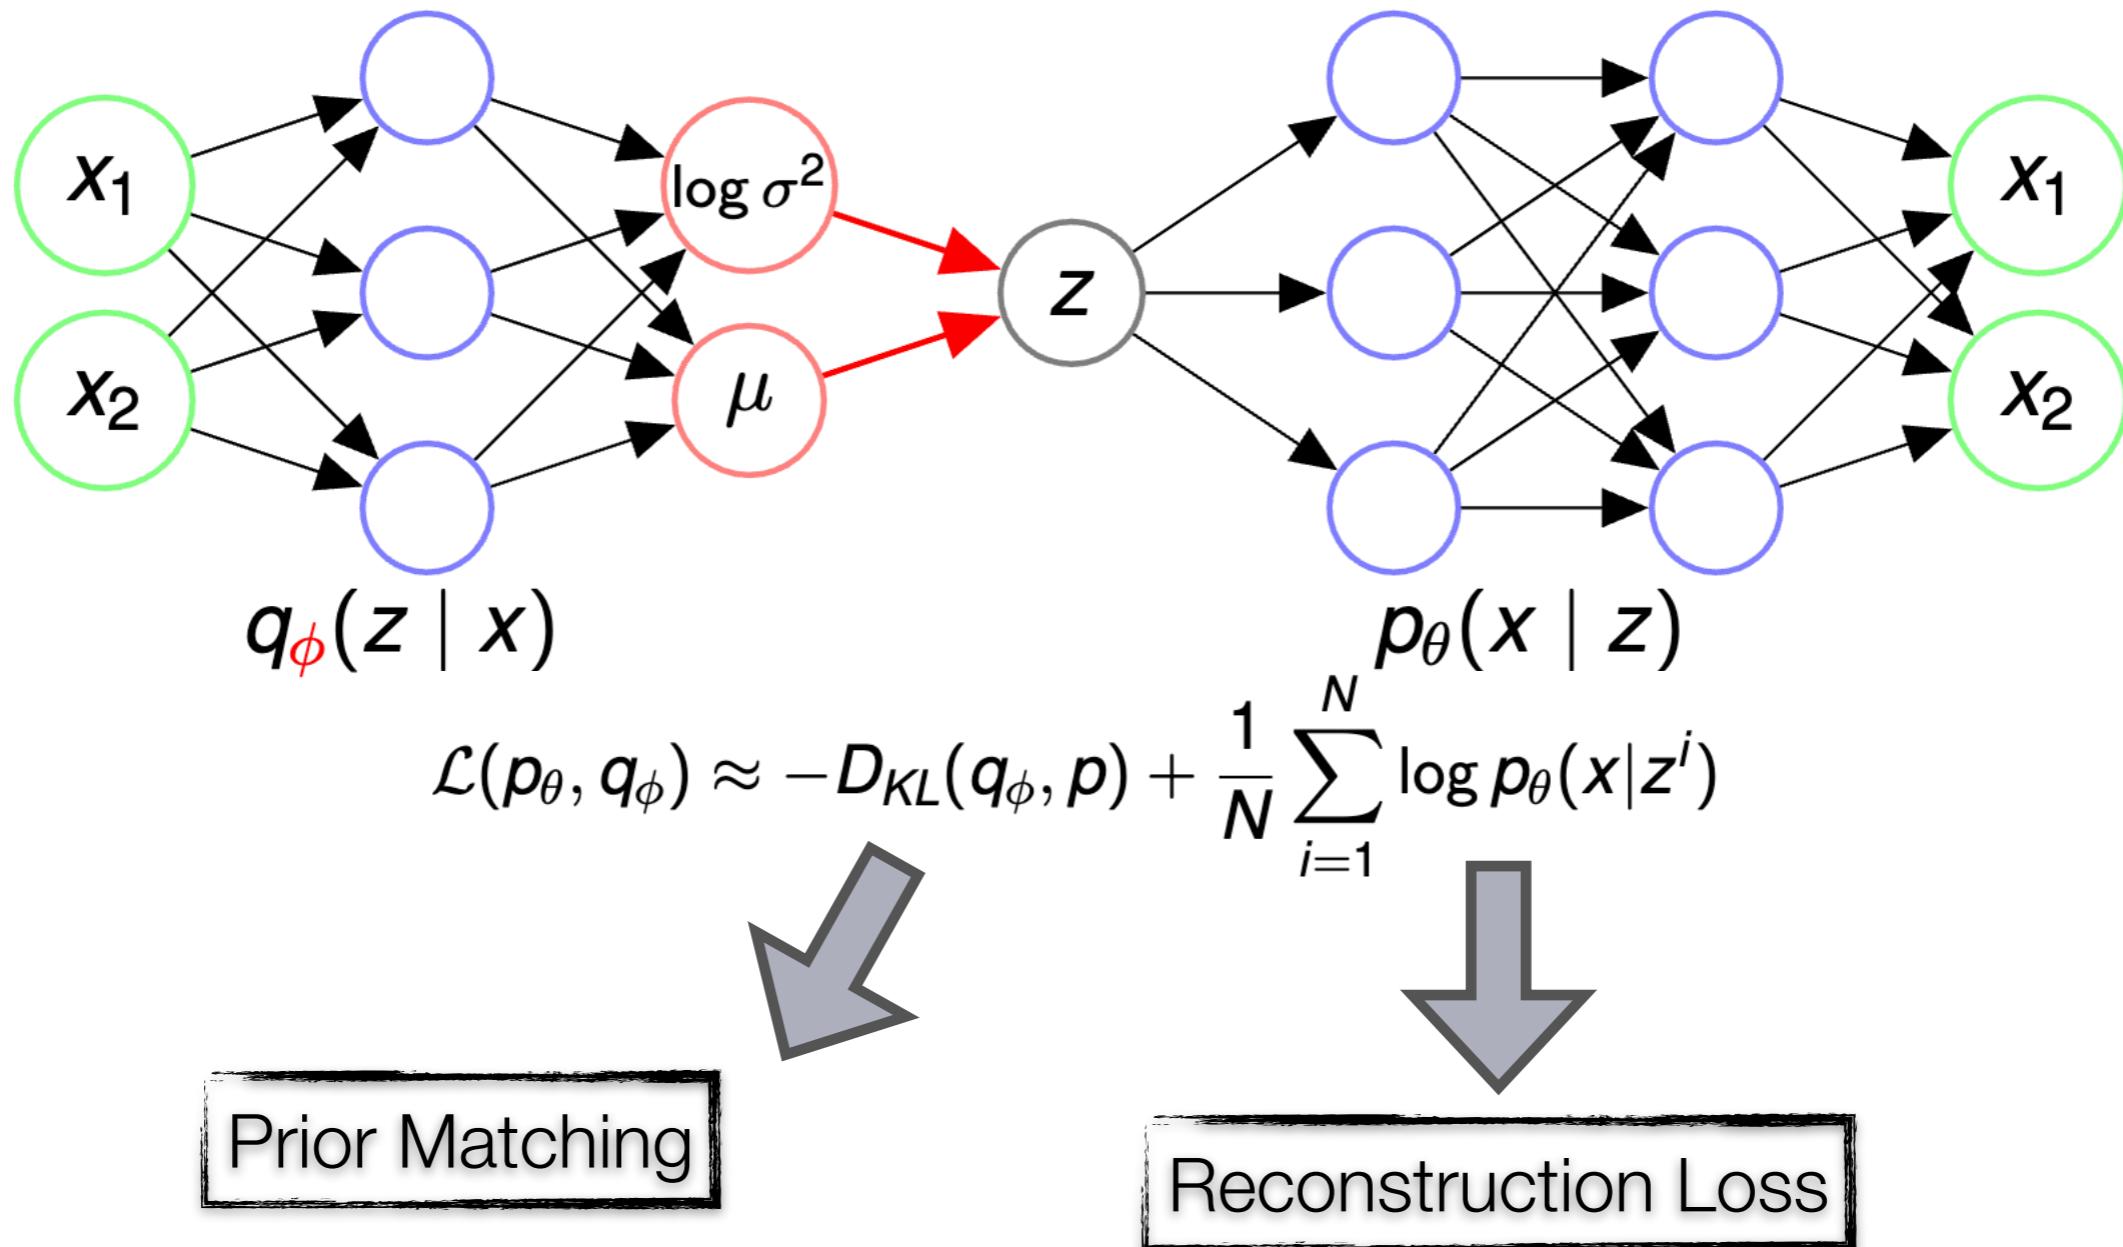
$$\mathcal{L}(p_\theta, q_\phi) \approx -D_{KL}(q_\phi, p) + \frac{1}{N} \sum_{i=1}^N \log p_\theta(x|z^i)$$

Notation:

- Encoder:  $q_\phi$  : mapping each input  $x$  to a Gaussian distribution  
 $\mathcal{N}(z | \mu_\phi(x), \sigma_\phi^2(x))$
- Decoder:  $p_\theta$ : mapping each hidden code  $z$  to a Gaussian distribution  
 $\mathcal{N}(x | \mu_\theta(z), \sigma^2 I)$

# Recap: Variational Auto-encoder

Recall from Lecture 20 we introduced the variational auto-encoder (VAE):



Note:  $p$  is a pre-defined simple prior distribution over the hidden codes  $\mathcal{Z}$ , usually chosen to be  $\mathcal{N}(0, I)$

# Recap: Variational Auto-encoder

Recall from Lecture 20 we introduced the variational auto-encoder (VAE):

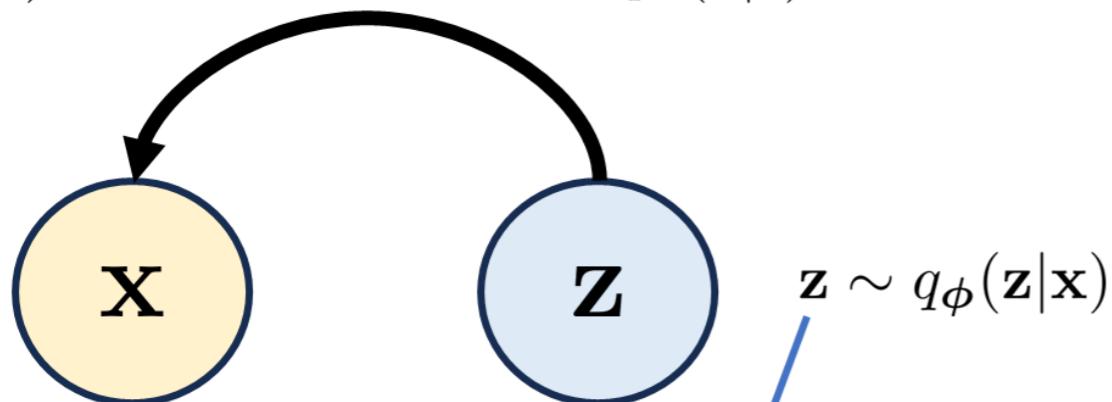
$$\mathcal{L}(p_\theta, q_\phi) \approx -D_{KL}(q_\phi, p) + \frac{1}{N} \sum_{i=1}^N \log p_\theta(x|z^i)$$

Prior Matching

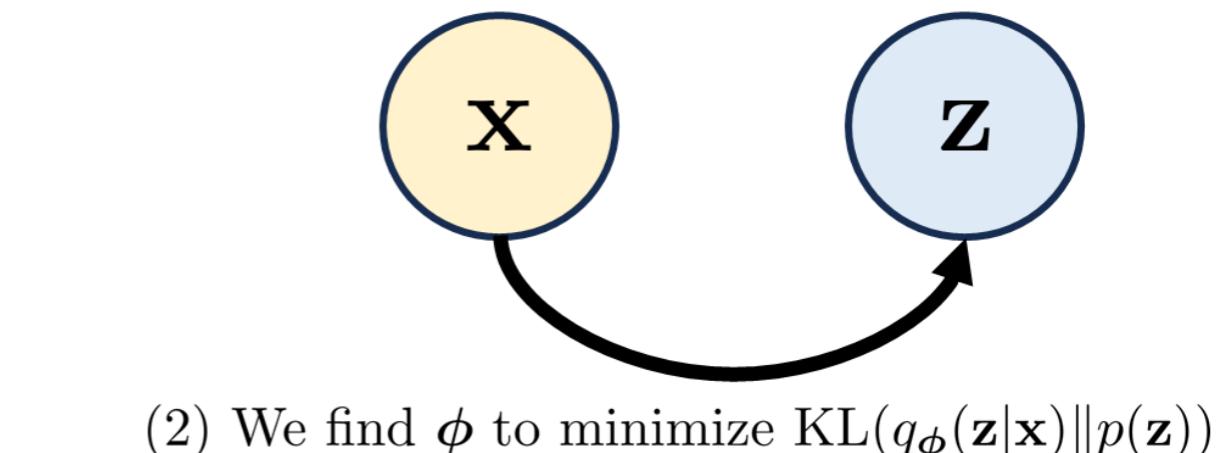


Reconstruction Loss

(2) We find  $\theta$  to maximize  $p_\theta(x|z)$



(1) During training, you tell us  $(z, x)$



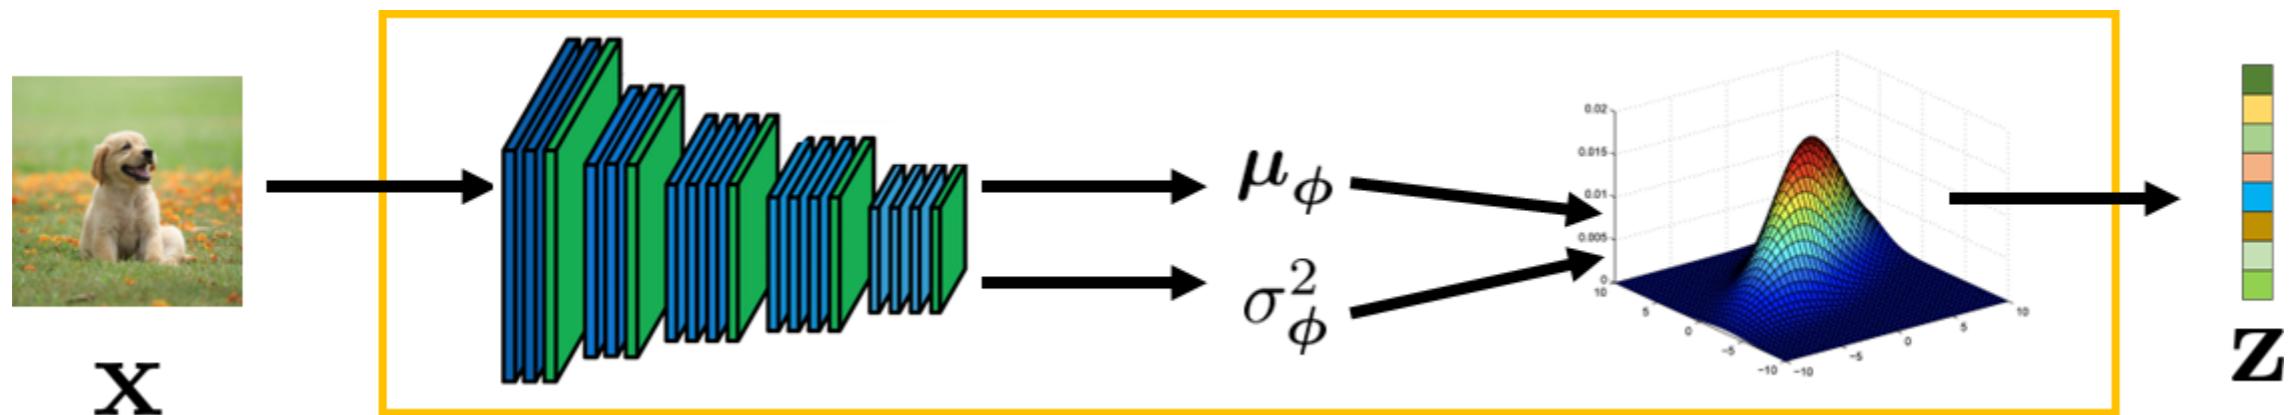
**Reconstruction**

**Prior Matching**

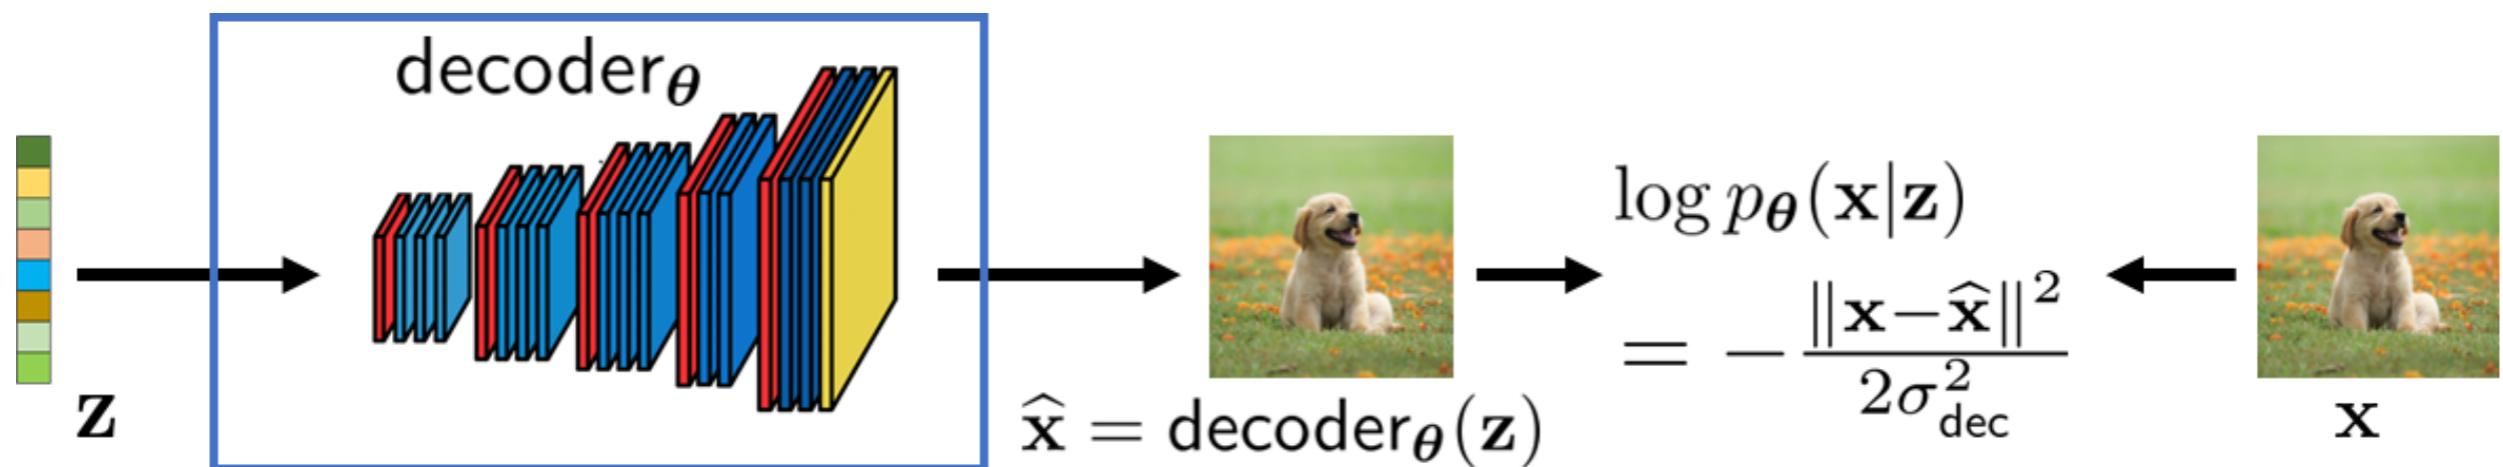
# Recap: Variational Auto-encoder

Training stage of VAE:

Encoder:



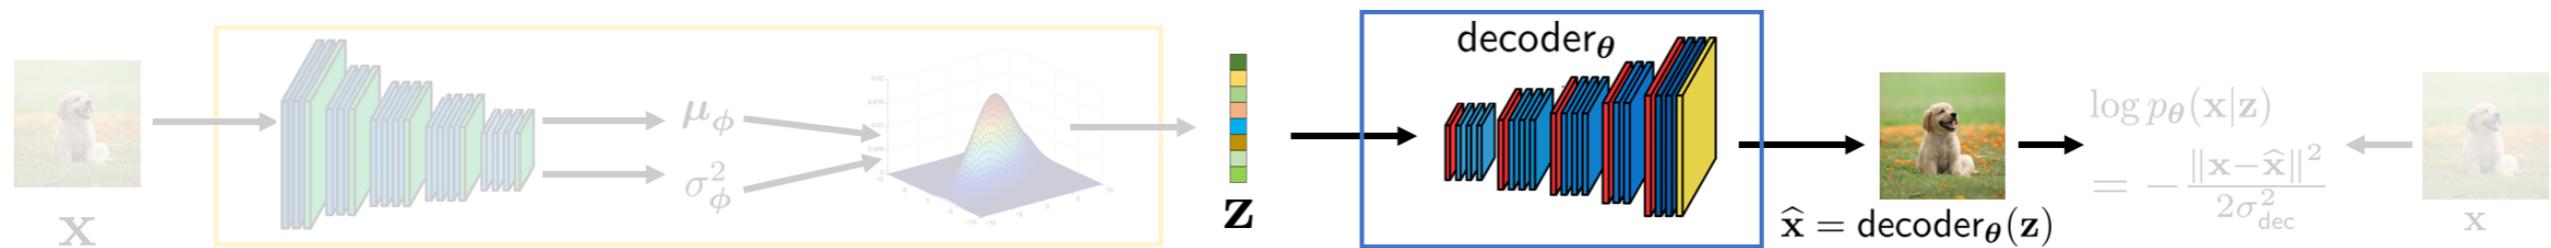
Decoder:



# Recap: Variational Auto-encoder

Inference stage of VAE:

Encoder gets dropped, only using the decoder

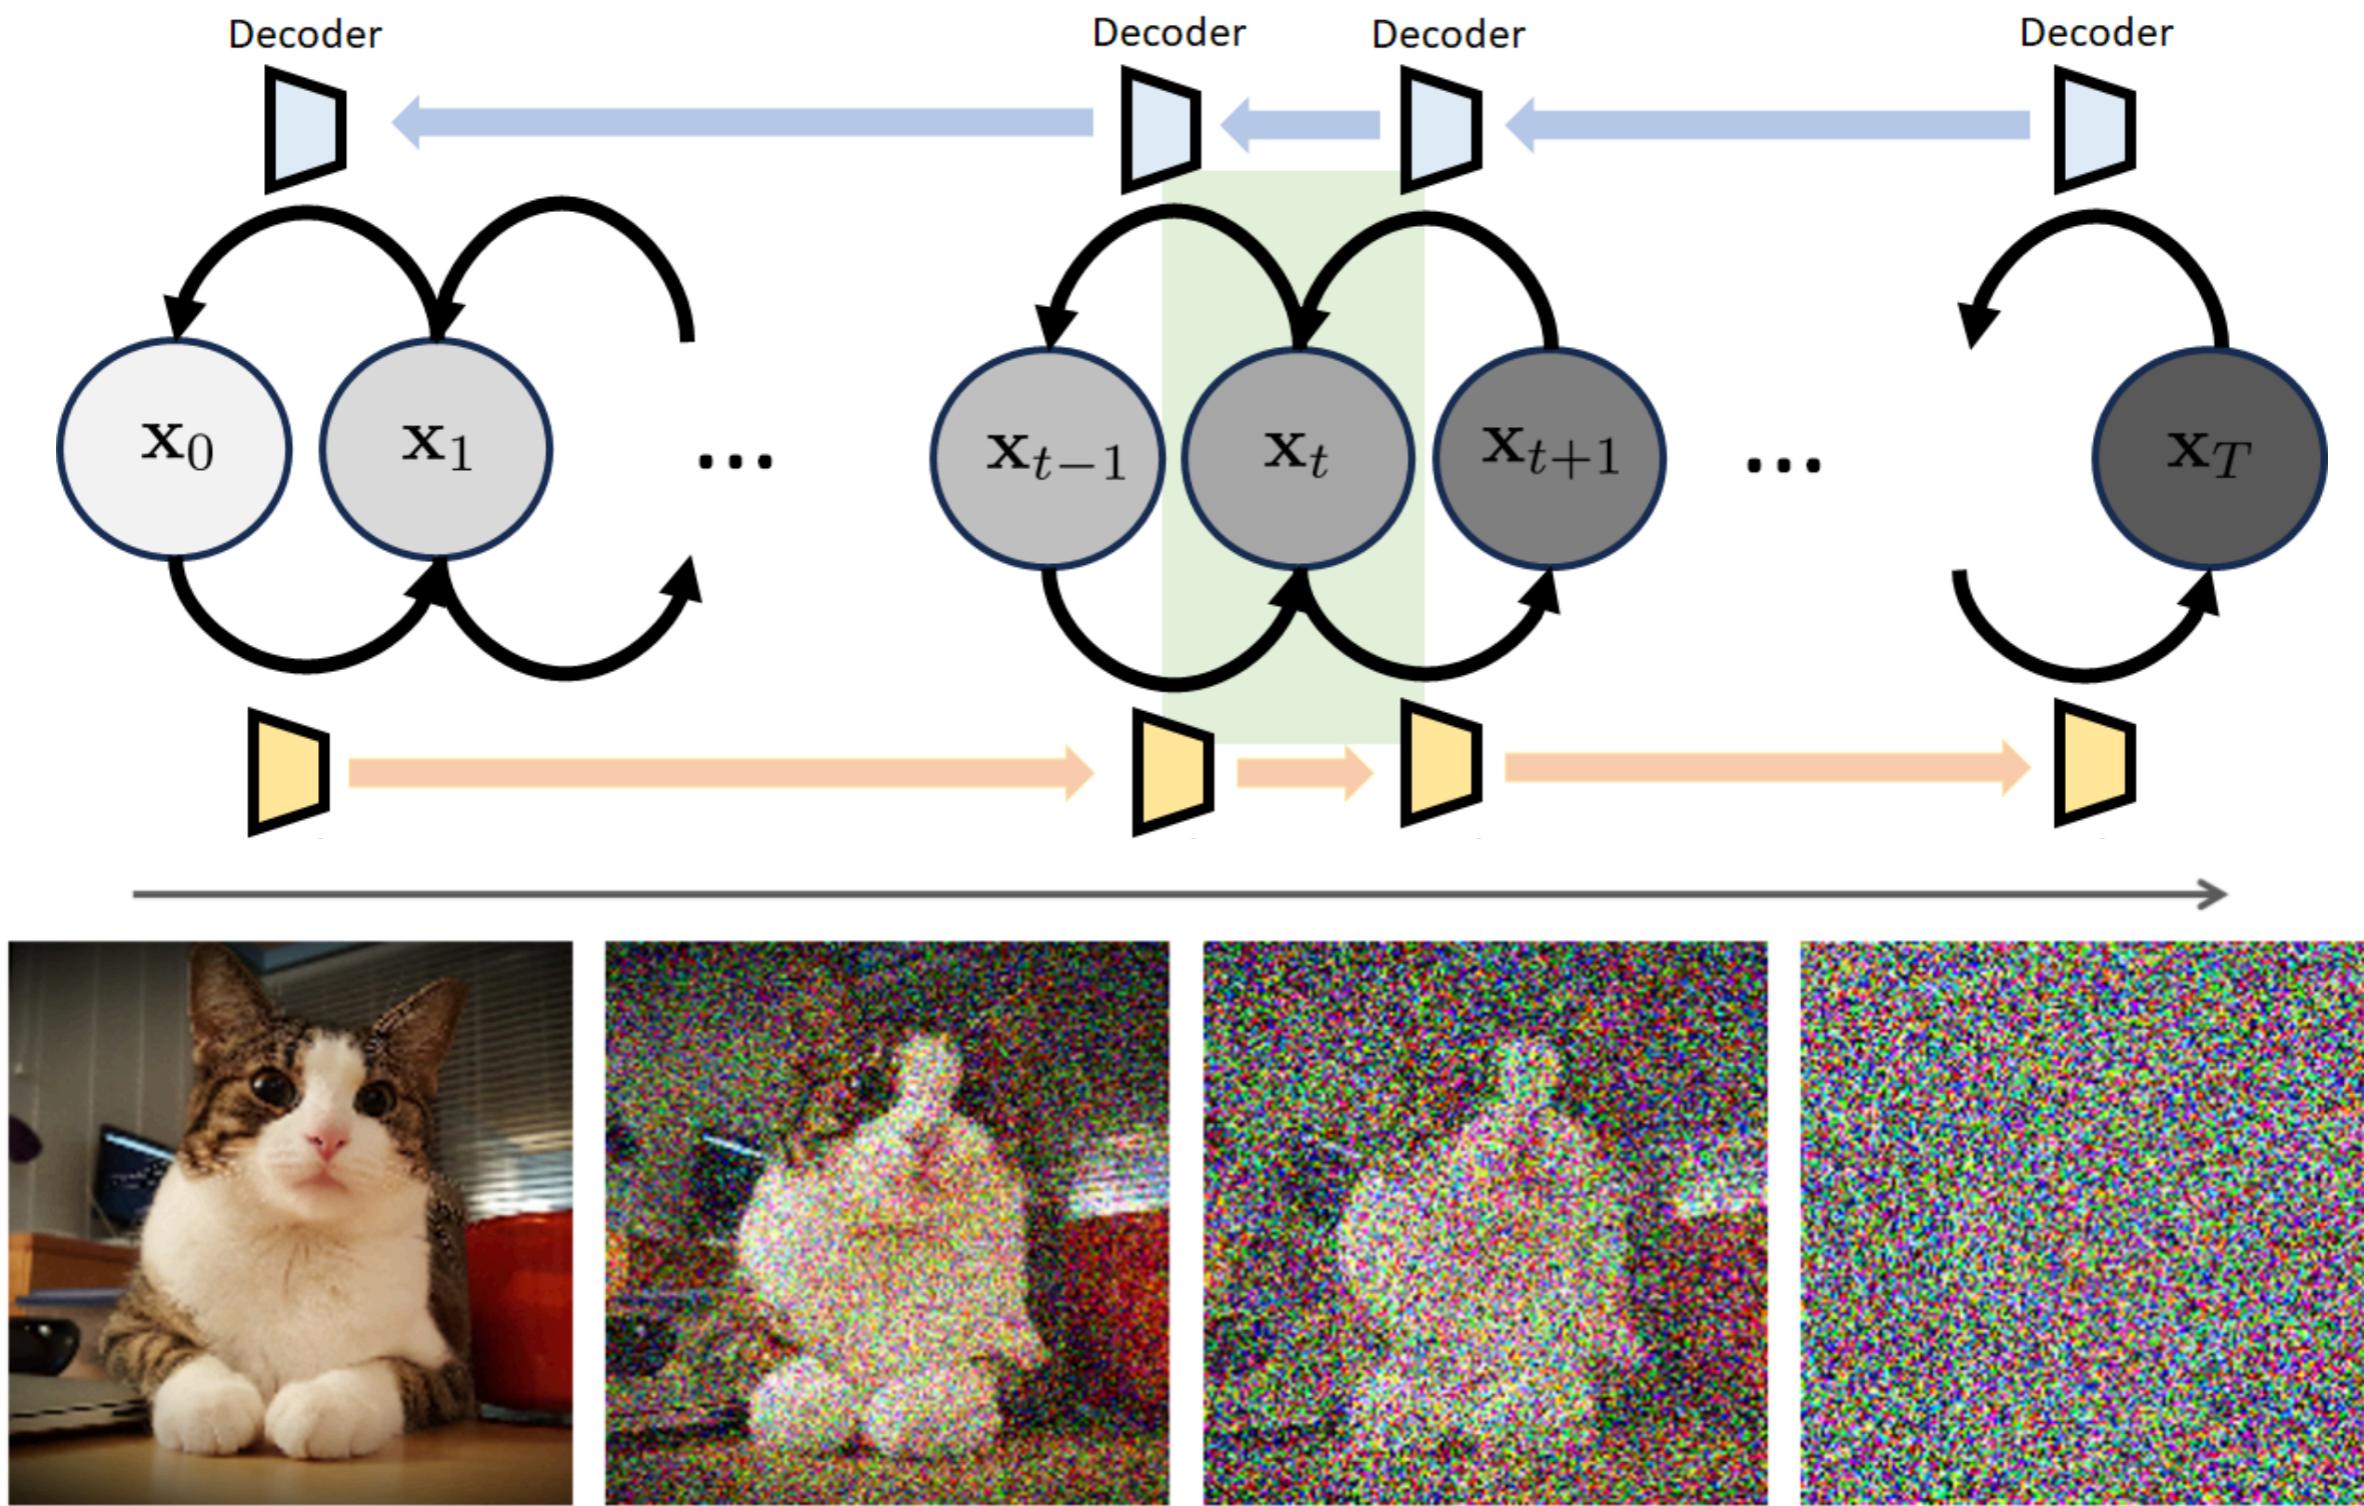


Like sampling from a GAN:

- First sample  $z \sim p = \mathcal{N}(0, I)$
- Pass the sampled  $z$  through the decoder to obtain  $\hat{x} = \text{Mean}(p_\theta(z))$

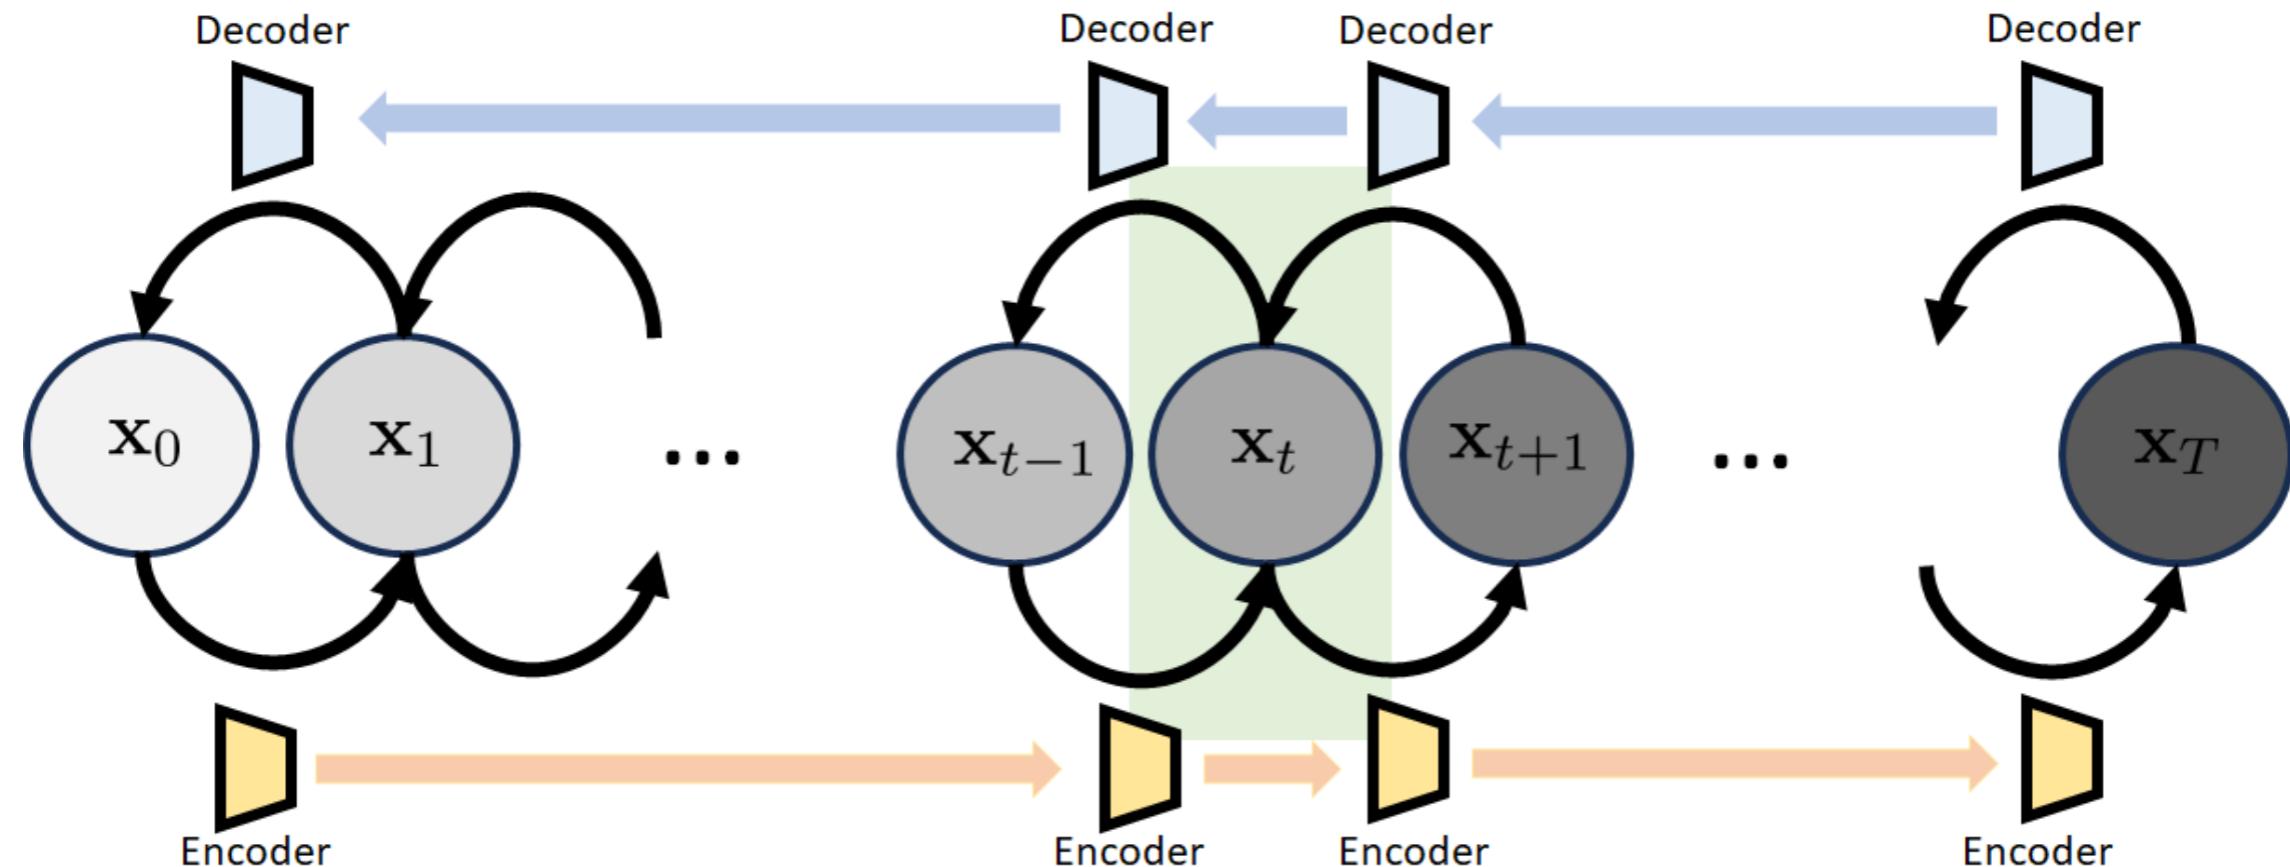
# Diffusion Models

Roughly speaking, diffusion models are iterative VAEs



# Diffusion Models

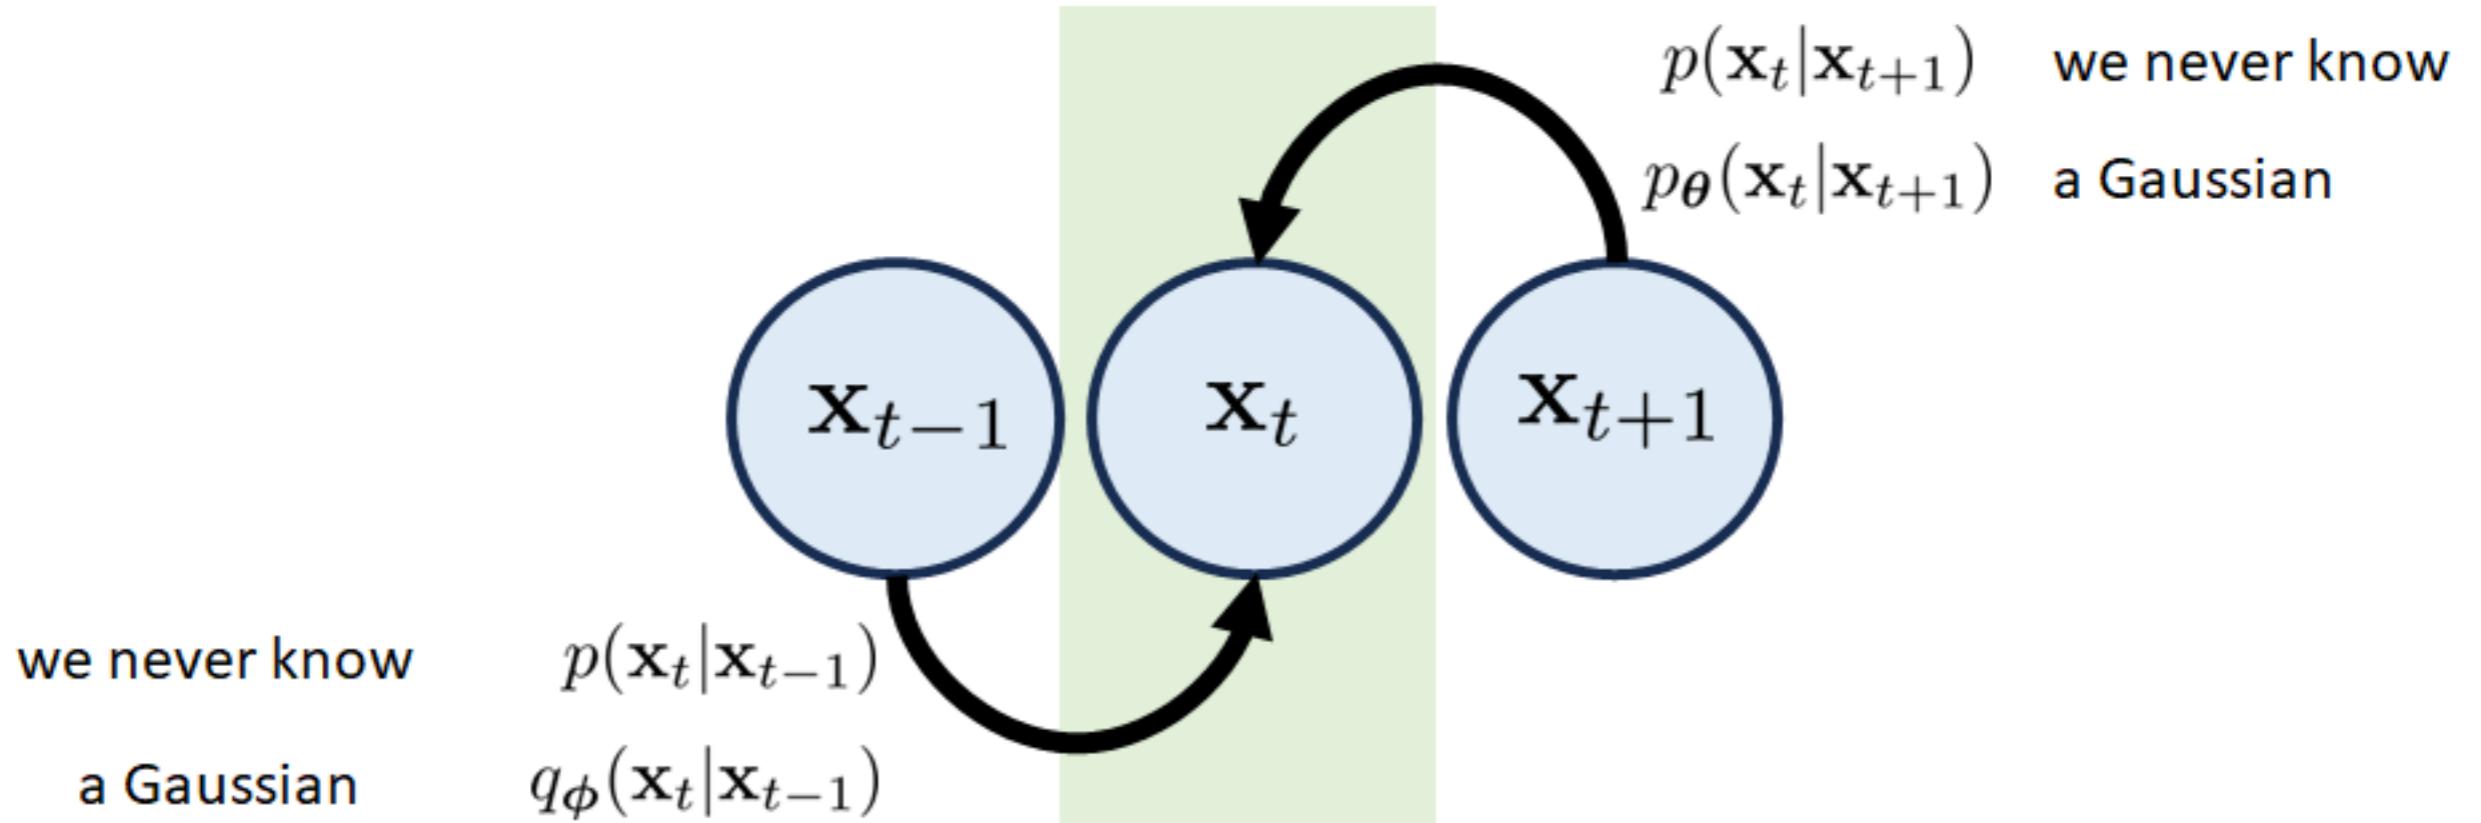
More formally,



- We have a sequence of states,  $x_0, \dots, x_T$
- There is an **unknown** underlying joint distribution  $p(x_0, \dots, x_T)$  over the sequence of states
- The initial state  $x_0 = x$  is the observable input image
- The end state  $x_T = z \sim \mathcal{N}(0, I)$  is supposed to be some hidden code sampled from a simple prior distribution, e.g., high-dim standard Gaussian

# Diffusion Models

Basic building block:

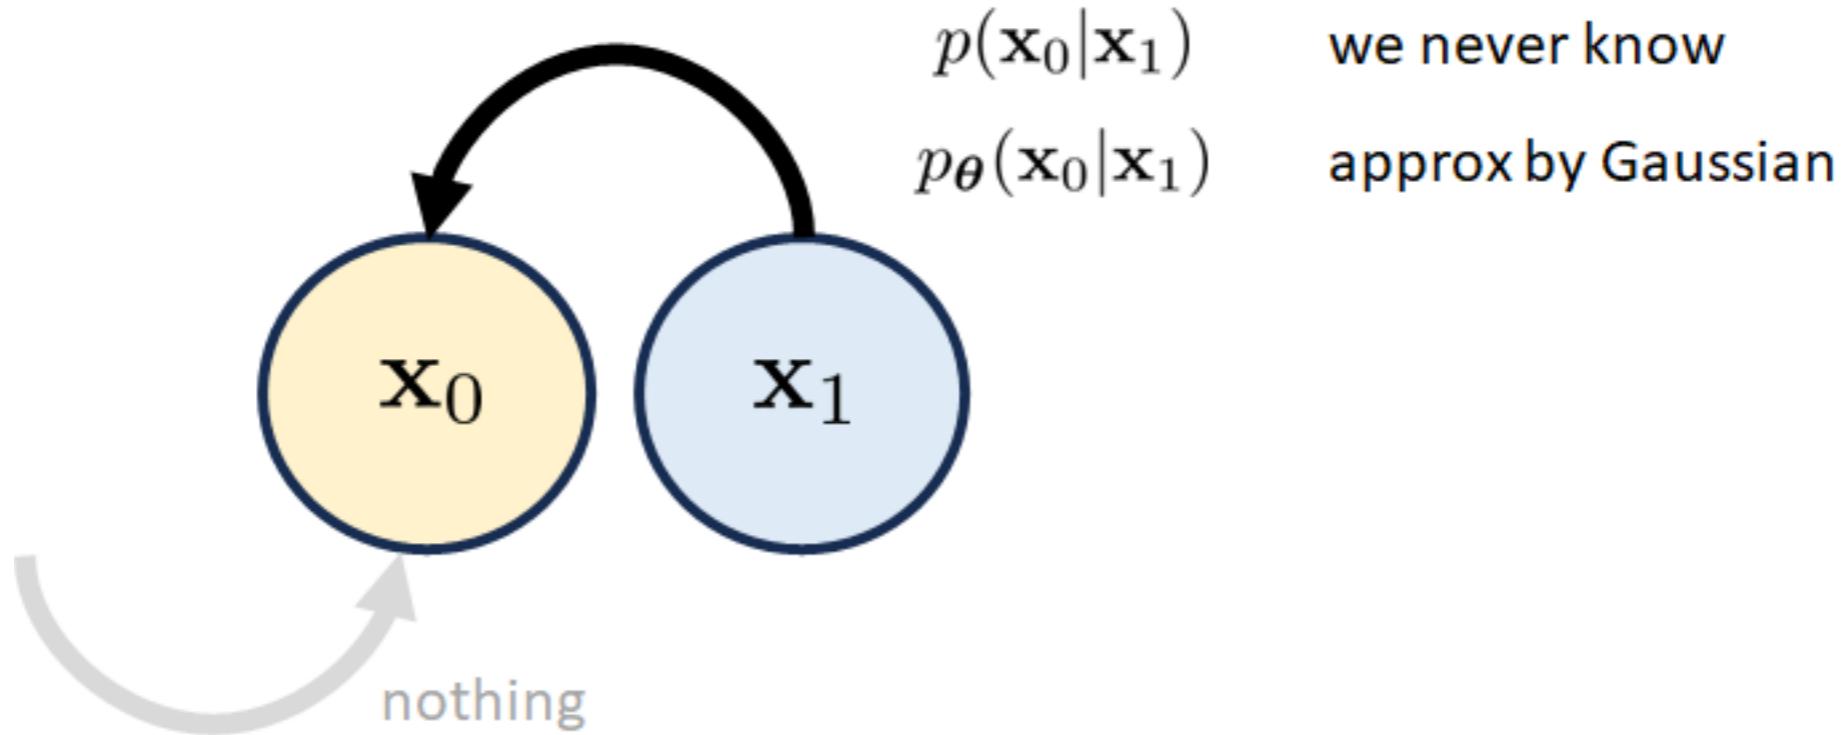


Similar to what we have in VAE:

- We use  $q_\phi$  as the **encoder** to approximate the underlying true conditional distribution  $p(x_t | x_{t-1})$
- We use  $p_\theta$  as the **decoder** to approximate the underlying true conditional distribution  $p(x_t | x_{t+1})$

# Diffusion Models

Initial block:

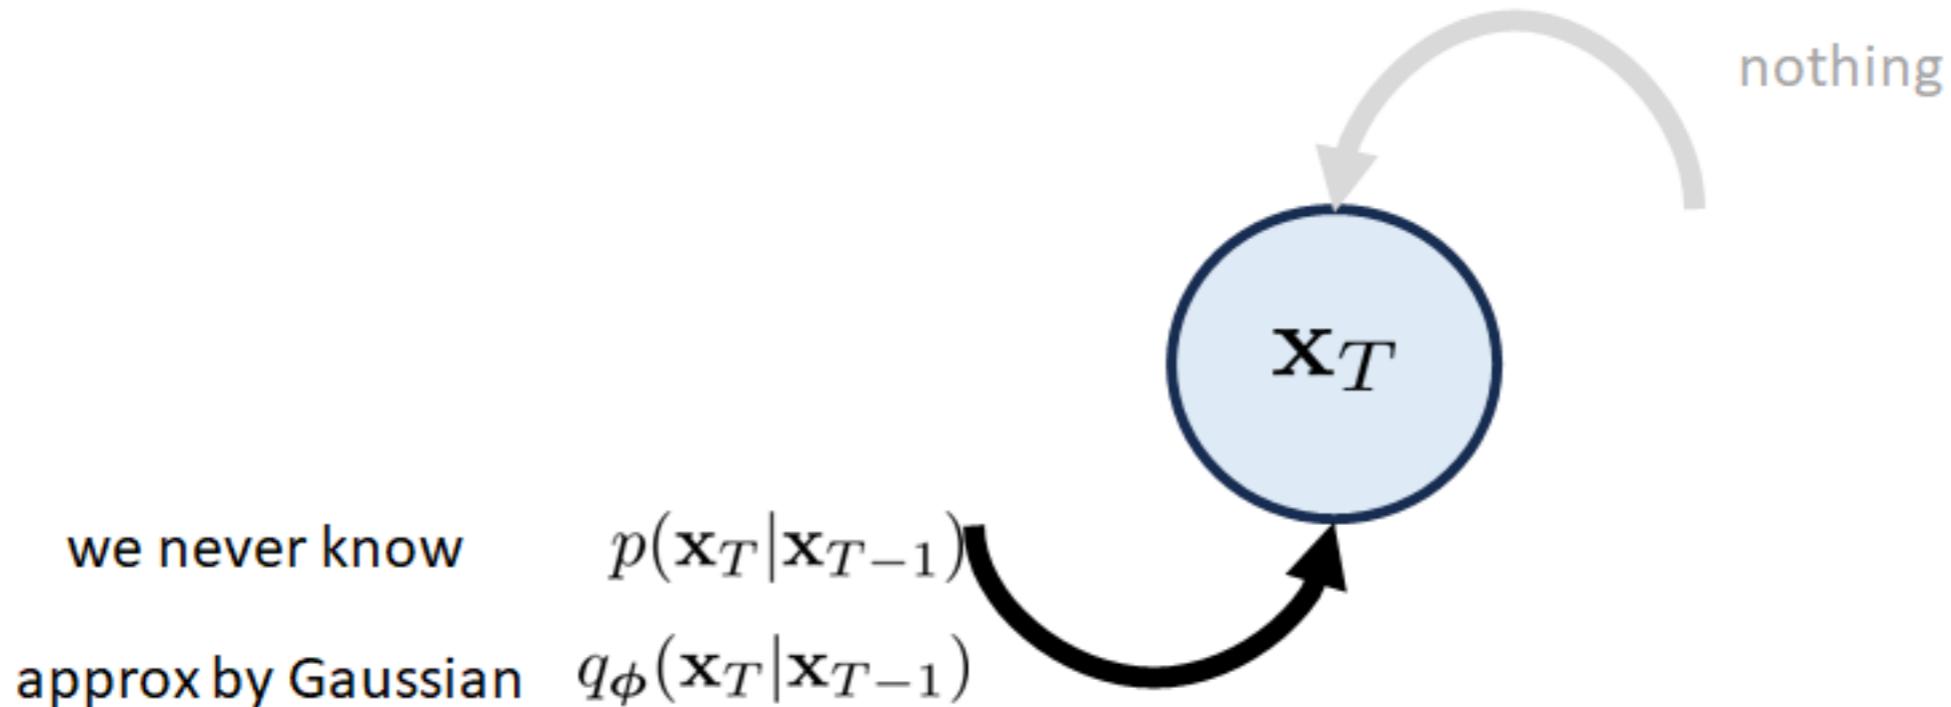


Similar to what we have in VAE:

- We use  $q_{\phi}$  as the **encoder** to approximate the underlying true conditional distribution  $p(x_t | x_{t-1})$
- We use  $p_{\theta}$  as the **decoder** to approximate the underlying true conditional distribution  $p(x_t | x_{t+1})$

# Diffusion Models

Final block:



Similar to what we have in VAE:

- We use  $q_\phi$  as the **encoder** to approximate the underlying true conditional distribution  $p(x_t | x_{t-1})$
- We use  $p_\theta$  as the **decoder** to approximate the underlying true conditional distribution  $p(x_t | x_{t+1})$

# Diffusion Models

---

Encoder  $q_\phi(x_t | x_{t-1})$ :

$$q_\phi(x_t | x_{t-1}) := \mathcal{N}\left(x_t | \sqrt{\alpha_t}x_{t-1}, (1 - \alpha_t)I\right)$$

The overall forward process (gradually adding noise to samples):

$$q_\phi(x_{1:T} | x_0) := \prod_{i=1}^T q_\phi(x_i | x_{i-1})$$

- The scaling factor  $0 < \alpha_t < 1$
- Equivalently, the above transition probability gives the following state equation

$$x_t = \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_t$$

where  $\epsilon_t \sim \mathcal{N}(0, I)$  is a random Gaussian noise with 0 mean and  $I$  covariance at the time step  $t \leq T$ .

Intuition: at each step, we add some Gaussian noise in hope that at the end when  $T$  is large, we will converge to  $\mathcal{N}(0, I)$ .

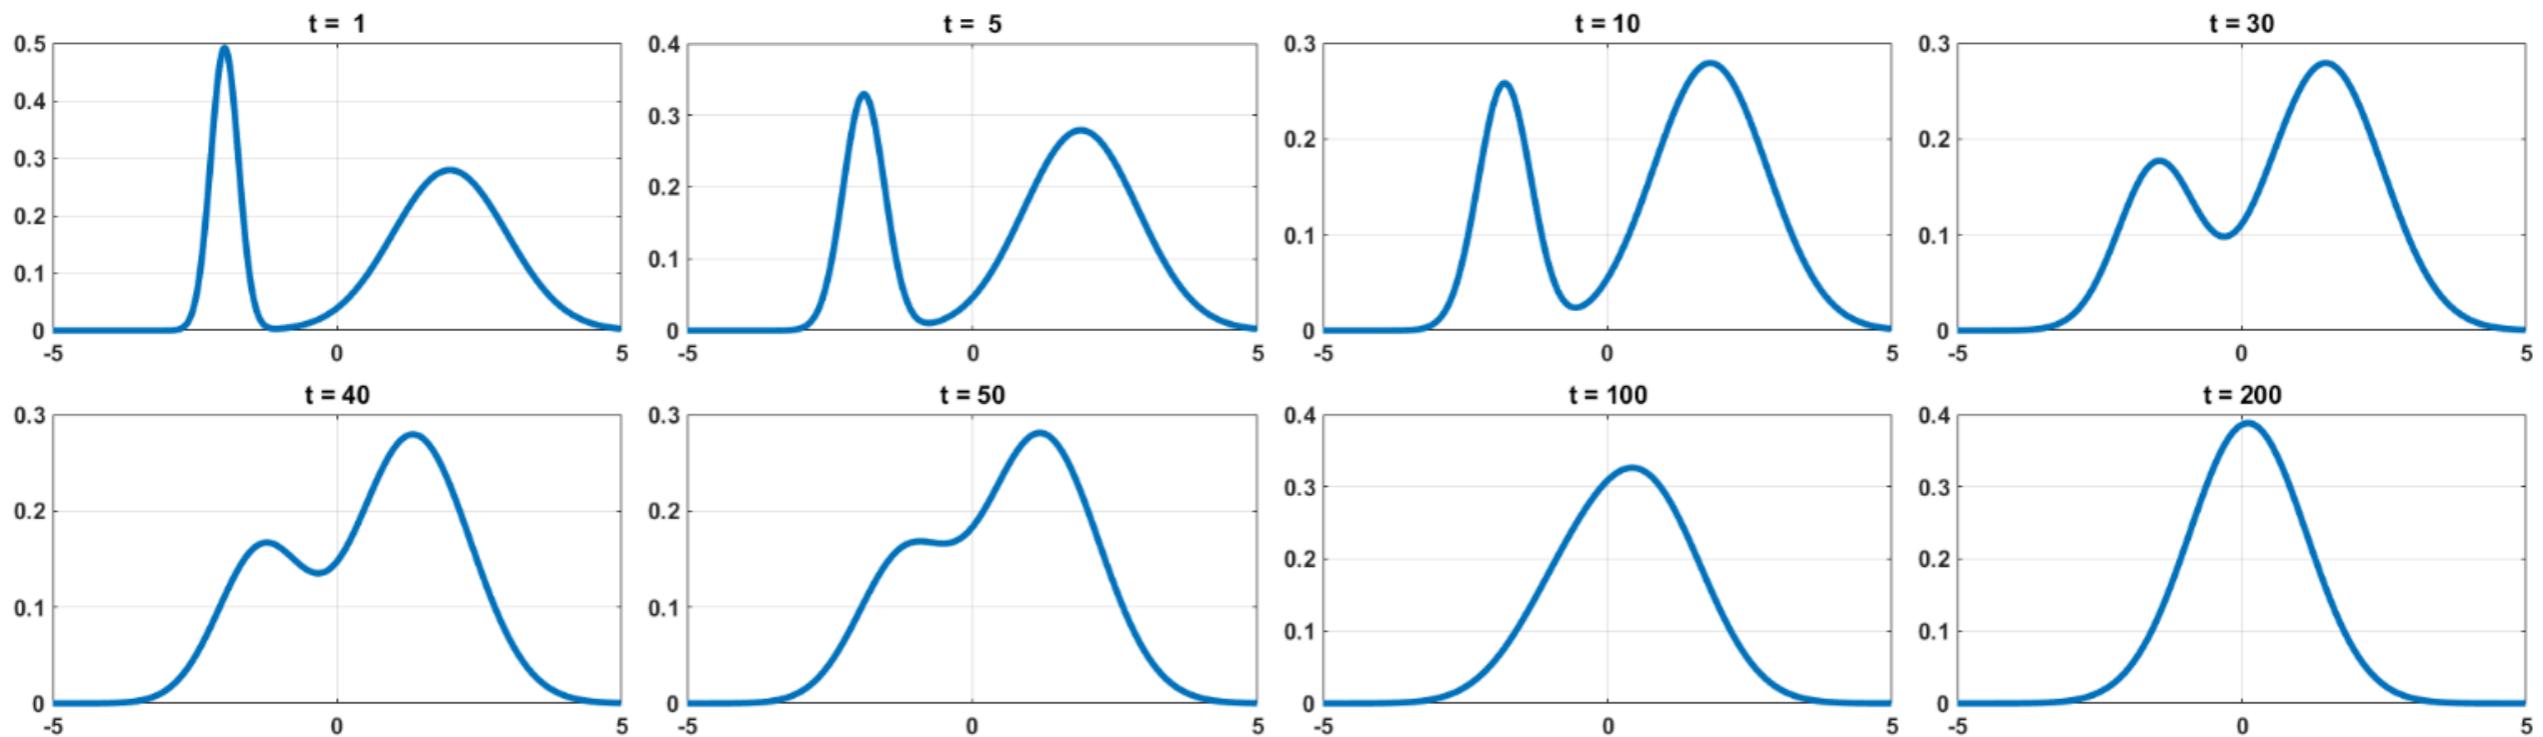
# Diffusion Models

Encoder  $q_\phi(x_t | x_{t-1})$ :

$$q_\phi(x_t | x_{t-1}) := \mathcal{N} \left( x_t | \sqrt{\alpha_t} x_{t-1}, (1 - \alpha_t) I \right)$$

Example:

$$x_0 \sim p_0(x) = \pi_1 \mathcal{N}(x | \mu_1, \sigma_1^2) + \pi_2 \mathcal{N}(x | \mu_2, \sigma_2^2)$$



# Diffusion Models

---

Encoder  $q_\phi(x_t | x_{t-1})$ :

$$q_\phi(x_t | x_{t-1}) := \mathcal{N}\left(x_t | \sqrt{\alpha_t}x_{t-1}, (1 - \alpha_t)I\right)$$

Why the specific constants  $\sqrt{\alpha_t}$  and  $1 - \alpha_t$ ?

Consider the special case where  $\alpha_t = \alpha \in (0,1), \forall t$ . By the iterative state transition equation, we have

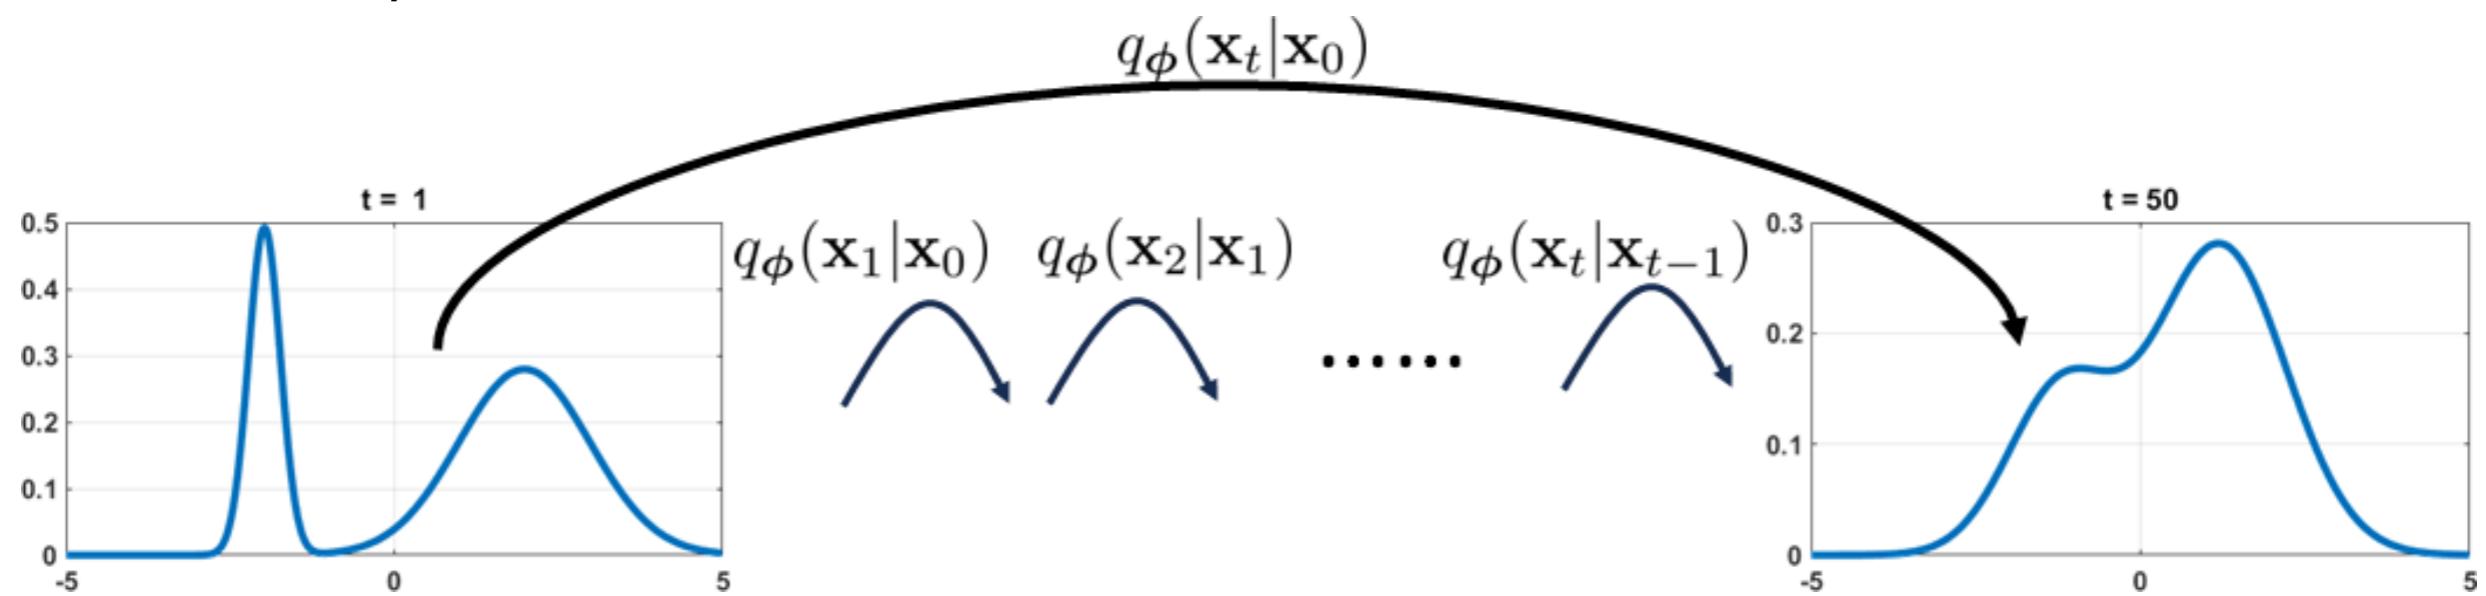
$$x_t \sim q_\phi(x_t | x_0) = \mathcal{N}(x_t | \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)$$

$$\text{where } \bar{\alpha}_t := \prod_{i=1}^t \alpha_i$$

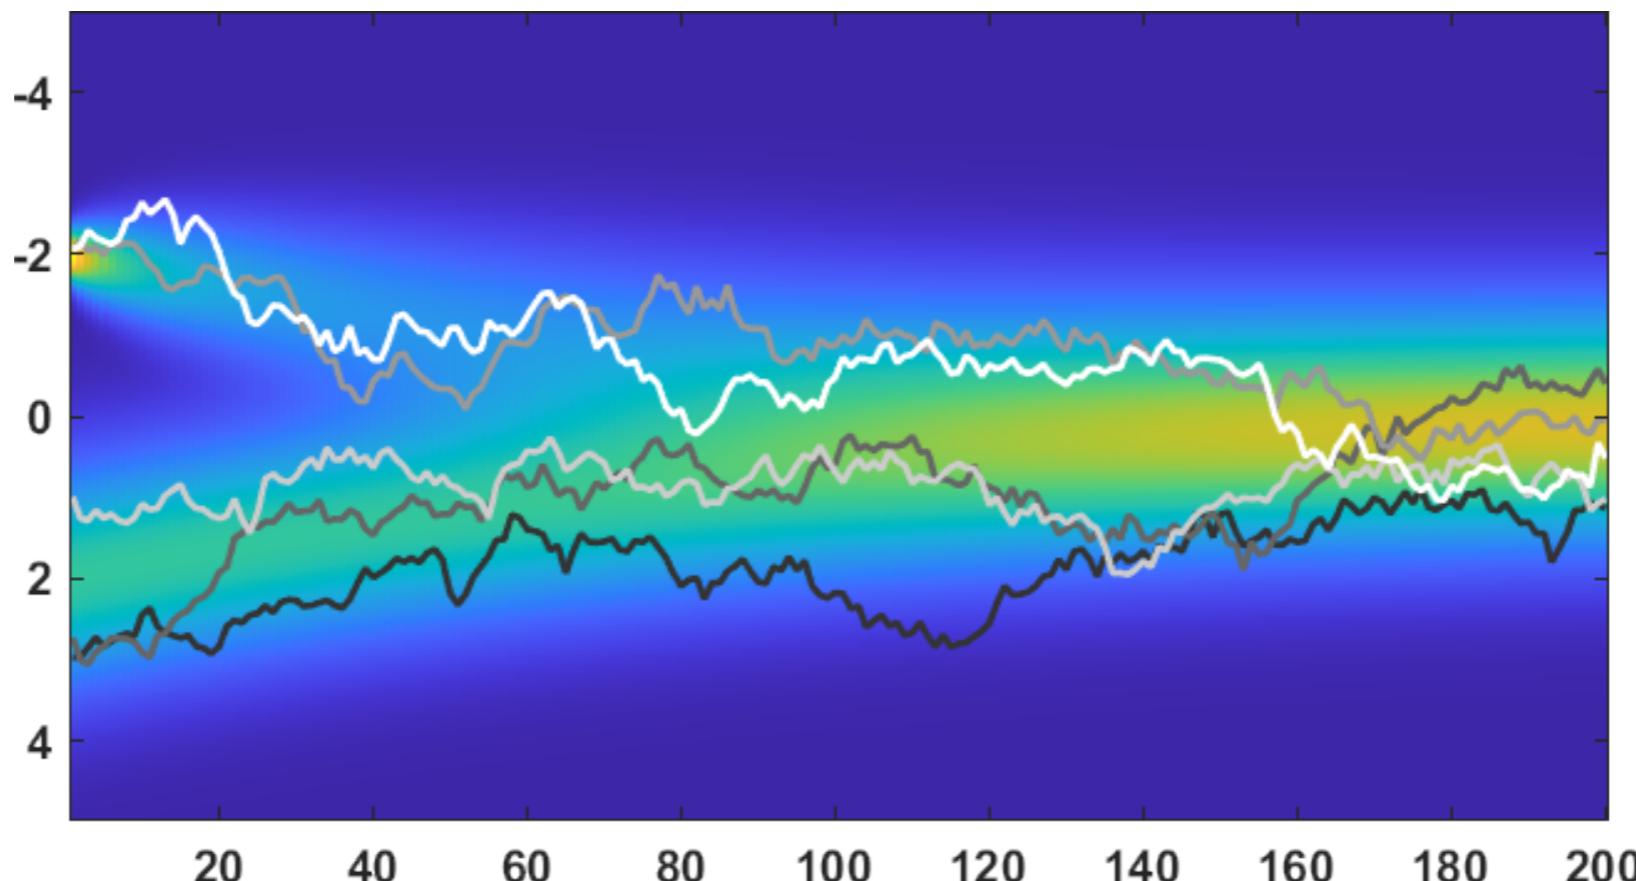
Convince yourself that as  $t \rightarrow \infty$ , the conditional distribution  $q_\phi(x_t | x_0)$  will converge to  $\mathcal{N}(0, I)$  for any initial distribution over  $x_0$ .

# Diffusion Models

Encoder  $q_\phi(x_t | x_{t-1})$ :



Trajectory of a Gaussian mixture turning into a standard Gaussian:



# Diffusion Models

Decoder  $p_\theta(x_t | x_{t+1})$ :

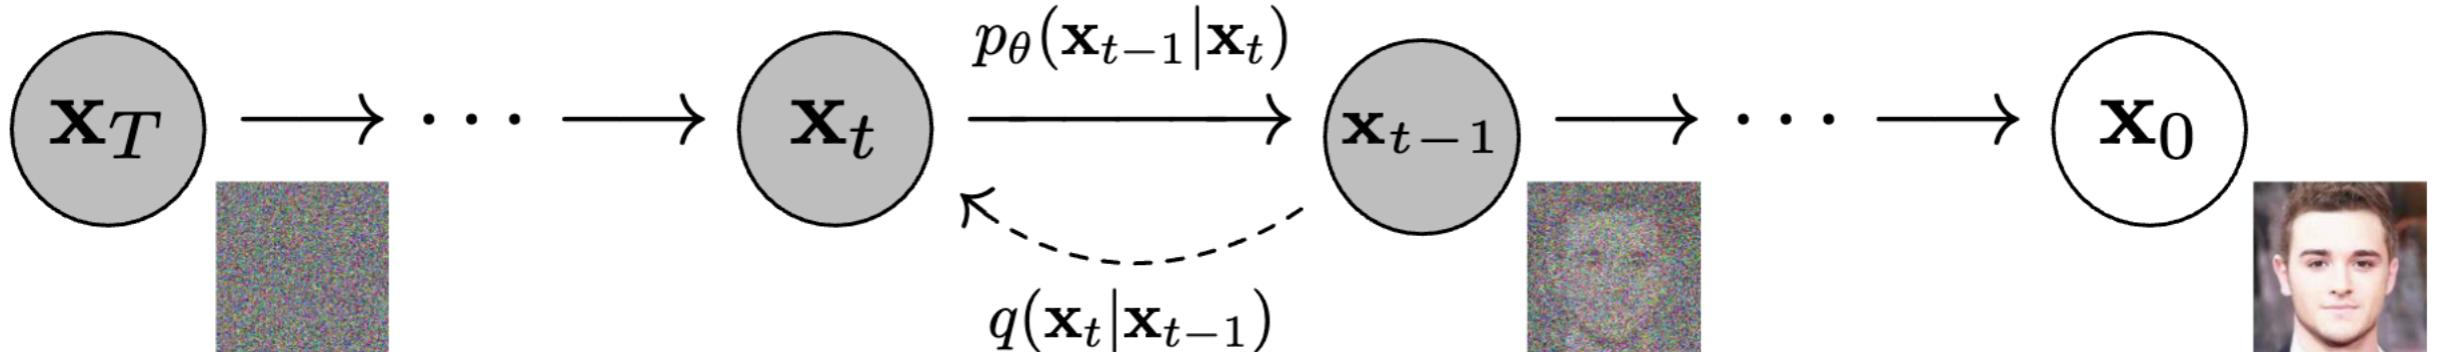
$$p_\theta(x_t | x_{t+1}) := \mathcal{N}(x_t | \mu_\theta(x_{t+1}, t+1), \sigma_\theta(x_{t+1}, t+1))$$

- Again, similar to the case of VAE, we use neural networks to learn the parameters of a Gaussian distributions, including both the mean and the (diagonal) covariance matrix

The overall backward process (gradually generating samples from noise):

$$p_\theta(x_{0:T}) := p(x_T) \prod_{i=1}^T p_\theta(x_{t-1} | x_t)$$

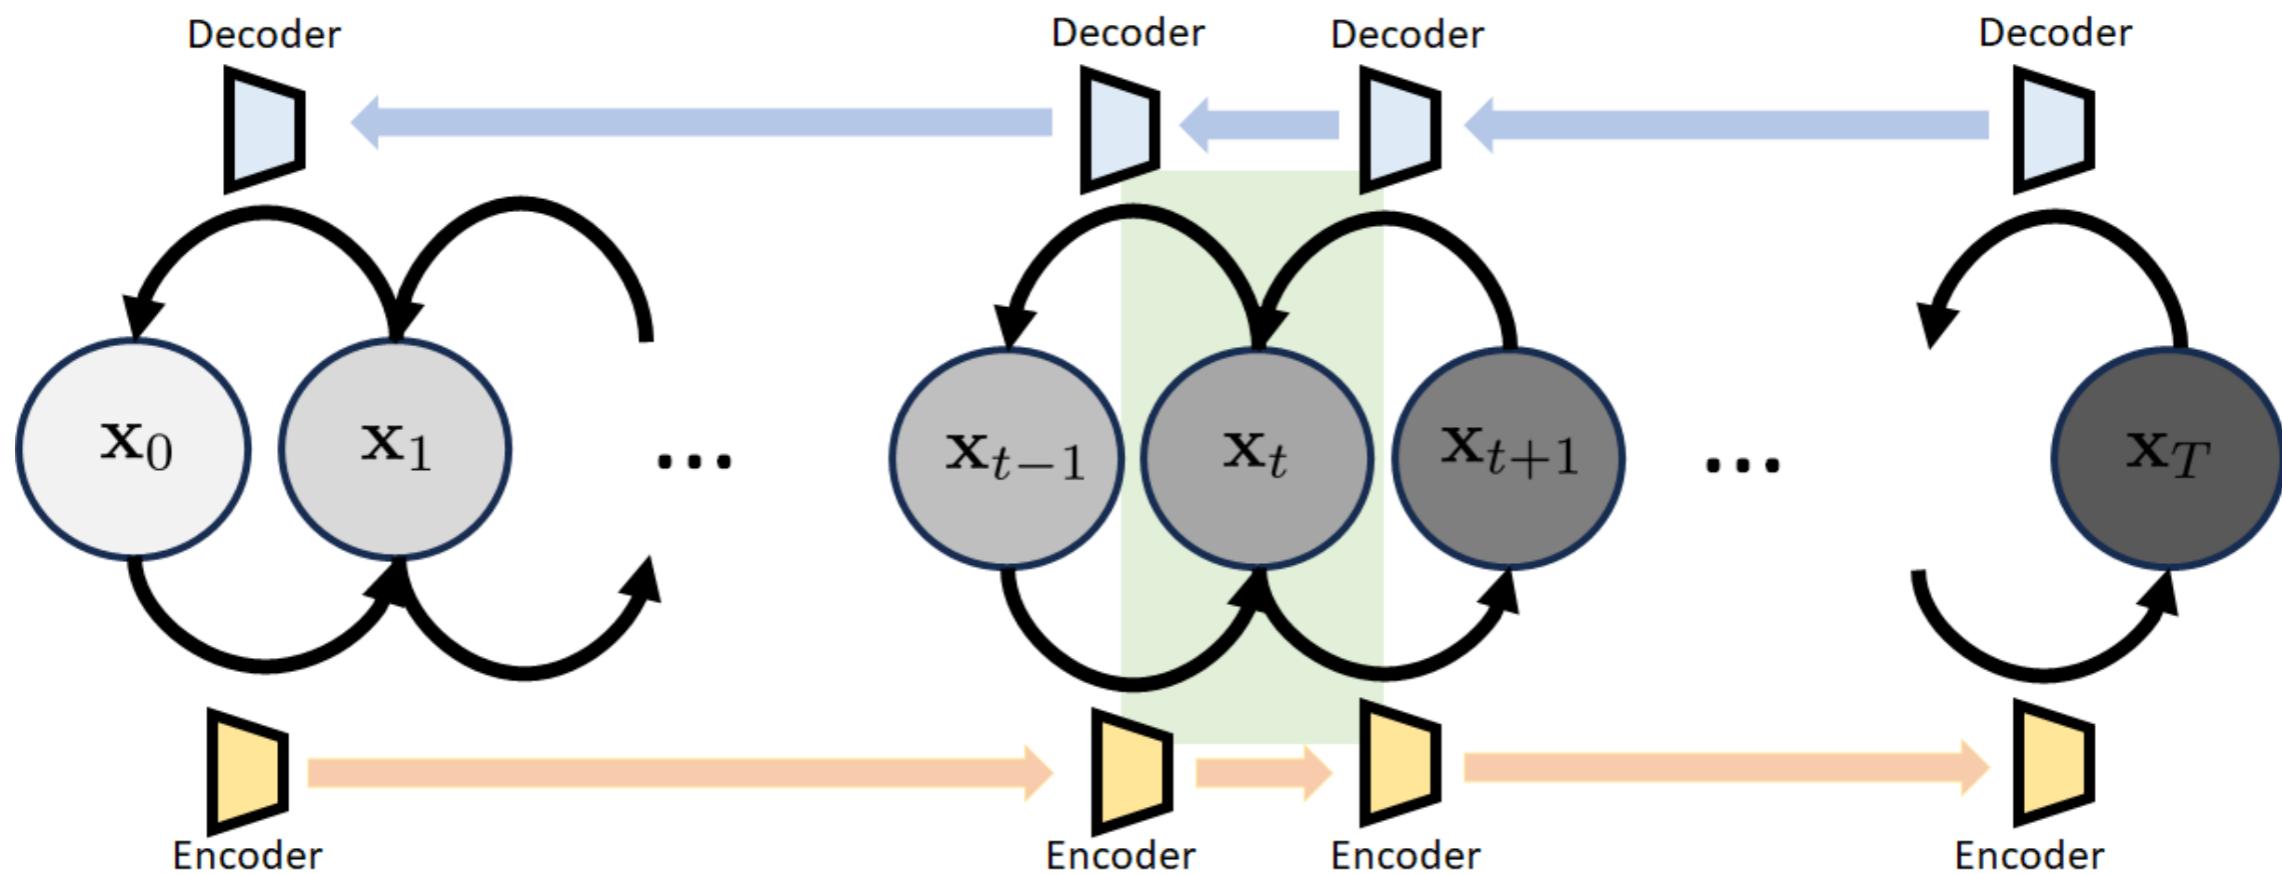
Note:  $p(x_T) = p(z) = \mathcal{N}(0, I)$



# Diffusion Models

## Training of diffusion models:

- Similar to VAE, we will need to derive the tractable evidence lower bound (ELBO) of the true log-likelihood function  $\log p_\theta(x_0)$
- Similar to VAE, the ELBO will include both the forward process (encoder) and the backward process (decoder)
- Due to its iterative nature, the overall loss will include more components than what we have in VAE



# Diffusion Models

---

Training of diffusion models:  
(HW 5)

$$\text{ELBO}_{\phi, \theta}(x) = \mathbb{E}_{q_\phi(x_1 | x_0)} [\log p_\theta(x_0 | x_1)]$$

$$- D_{\text{KL}}(q_\phi(x_T | x_0) \| p(x_T))$$

$$- \sum_{t=2}^T \mathbb{E}_{q_\phi(x_t | x_0)} \left[ D_{\text{KL}}(q_\phi(x_{t-1} | x_t, x_0) \| p_\theta(x_{t-1} | x_t)) \right]$$

- **(Reconstruction loss):** The first term measures how good the initial block is
- **(Prior matching):** The second term measures how good the final block is
- **(Consistency):** The last term measures how good the transition blocks are

# Diffusion Models

---

Training of diffusion models:

$$\begin{aligned} \text{ELBO}_{\phi, \theta}(x_0) &= \mathbb{E}_{q_\phi(x_1 | x_0)} [\log p_\theta(x_0 | x_1)] \\ &\quad - D_{\text{KL}}(q_\phi(x_T | x_0) \| p(x_T)) \\ &\quad - \sum_{t=2}^T \mathbb{E}_{q_\phi(x_t | x_0)} \left[ D_{\text{KL}}(q_\phi(x_{t-1} | x_t, x_0) \| p_\theta(x_{t-1} | x_t)) \right] \end{aligned}$$

(HW 5): How to compute  $q_\phi(x_{t-1} | x_t, x_0)$ ?

- Use the Bayes theorem
- It is also a Gaussian distribution (why?)

# Diffusion Models

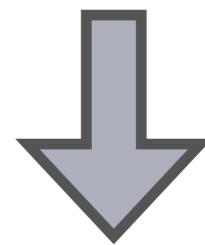
## Training of diffusion models:

After around ~5 pages of algebraic derivations (no conceptual challenges, though):

$$\text{ELBO}_\theta(x_0) = - \sum_{t=1}^T \frac{1}{2\sigma^2} \frac{(1 - \alpha_t)\bar{\alpha}_{t-1}}{(1 - \bar{\alpha}_t)^2} \mathbb{E}_{q_\phi(x_t | x_0)} [\|\hat{x}_\theta(x_t) - x_0\|_2^2]$$

Forward process

Reconstruction loss



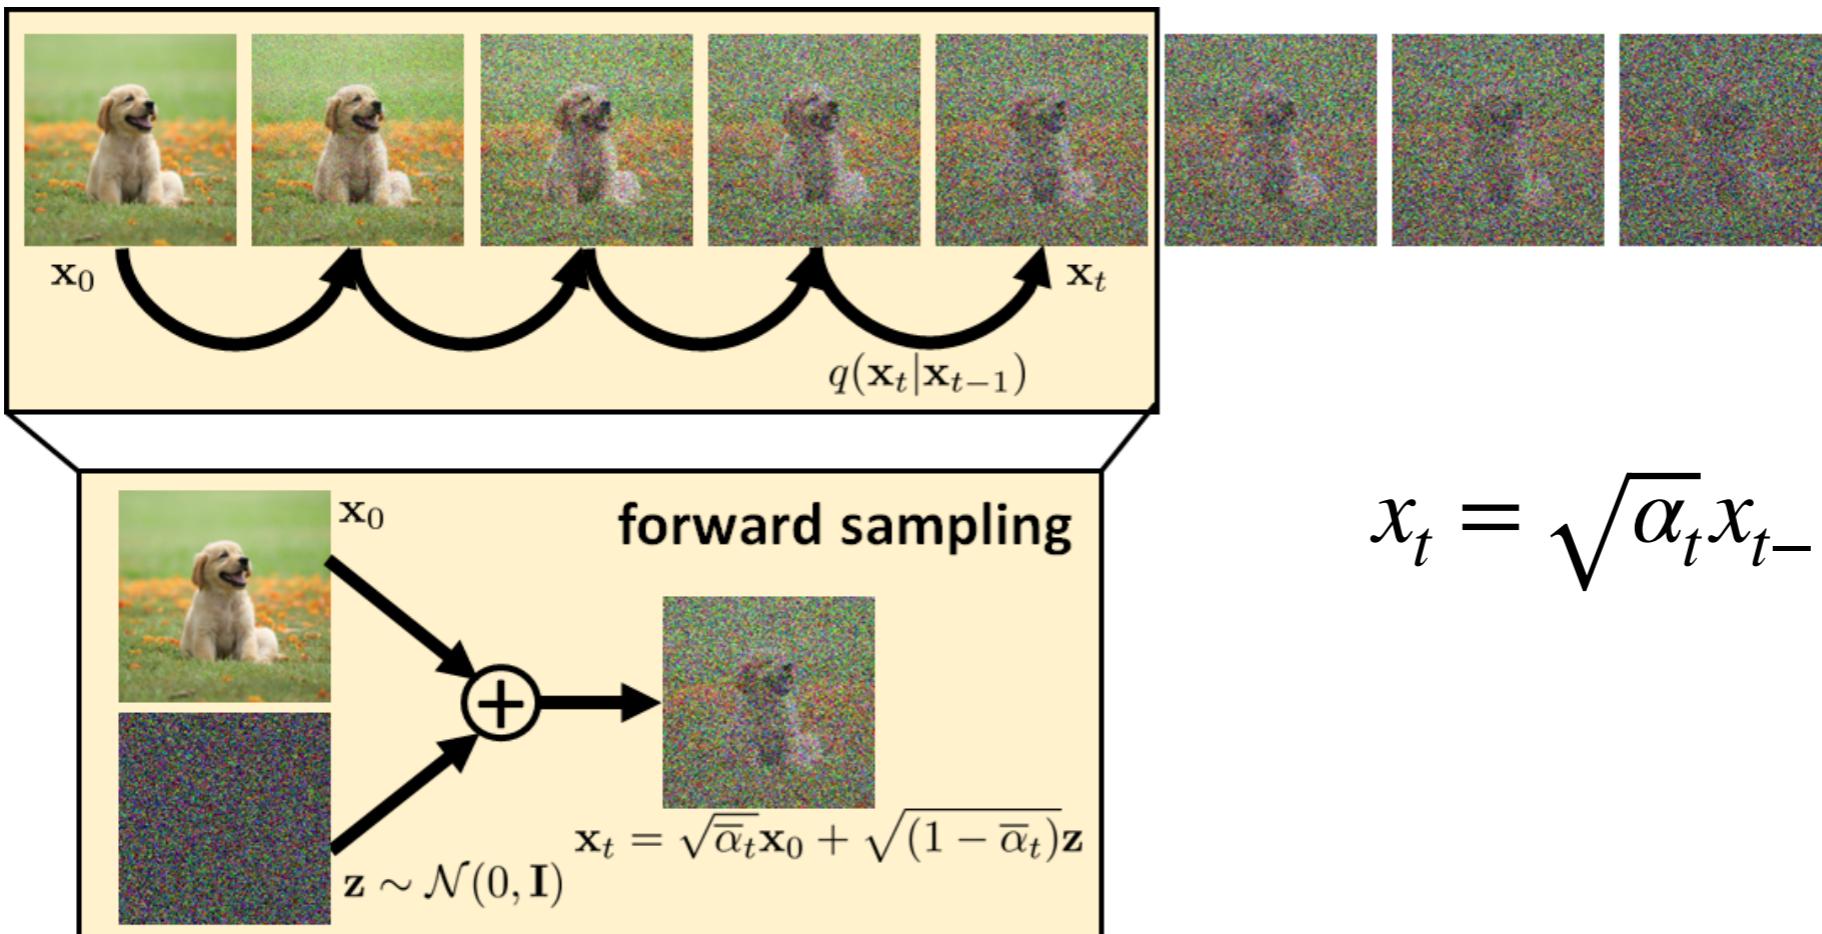
# Diffusion Models

## Training of diffusion models:

After around ~5 pages of algebraic derivations (no conceptual challenges, though):

$$\text{ELBO}_\theta(x_0) = - \sum_{t=1}^T \frac{1}{2\sigma^2} \frac{(1 - \alpha_t)\bar{\alpha}_{t-1}}{(1 - \bar{\alpha}_t)^2} \mathbb{E}_{q_\phi(x_t | x_0)} [\|\hat{x}_\theta(x_t) - x_0\|_2^2]$$

- The overall loss function is, again, reconstruction from noisy samples



# Diffusion Models

---

## Training of diffusion models:

After around ~5 pages of algebraic derivations (no conceptual challenges, though):

$$\text{ELBO}_\theta(x_0) = - \sum_{t=1}^T \frac{1}{2\sigma^2} \frac{(1 - \alpha_t)\bar{\alpha}_{t-1}}{(1 - \bar{\alpha}_t)^2} \mathbb{E}_{q_\phi(x_t | x_0)} [\|\hat{x}_\theta(x_t) - x_0\|_2^2]$$

- The overall loss function is, again, reconstruction from noisy samples
- Our parametrized model (neural network) is now  $\hat{x}_\theta : \mathcal{X} \rightarrow \mathcal{X}$ , which takes a noise image and gives back a denoised one
- The denoising process is carefully scheduled along the whole sequential process  $t = 1, \dots, T$

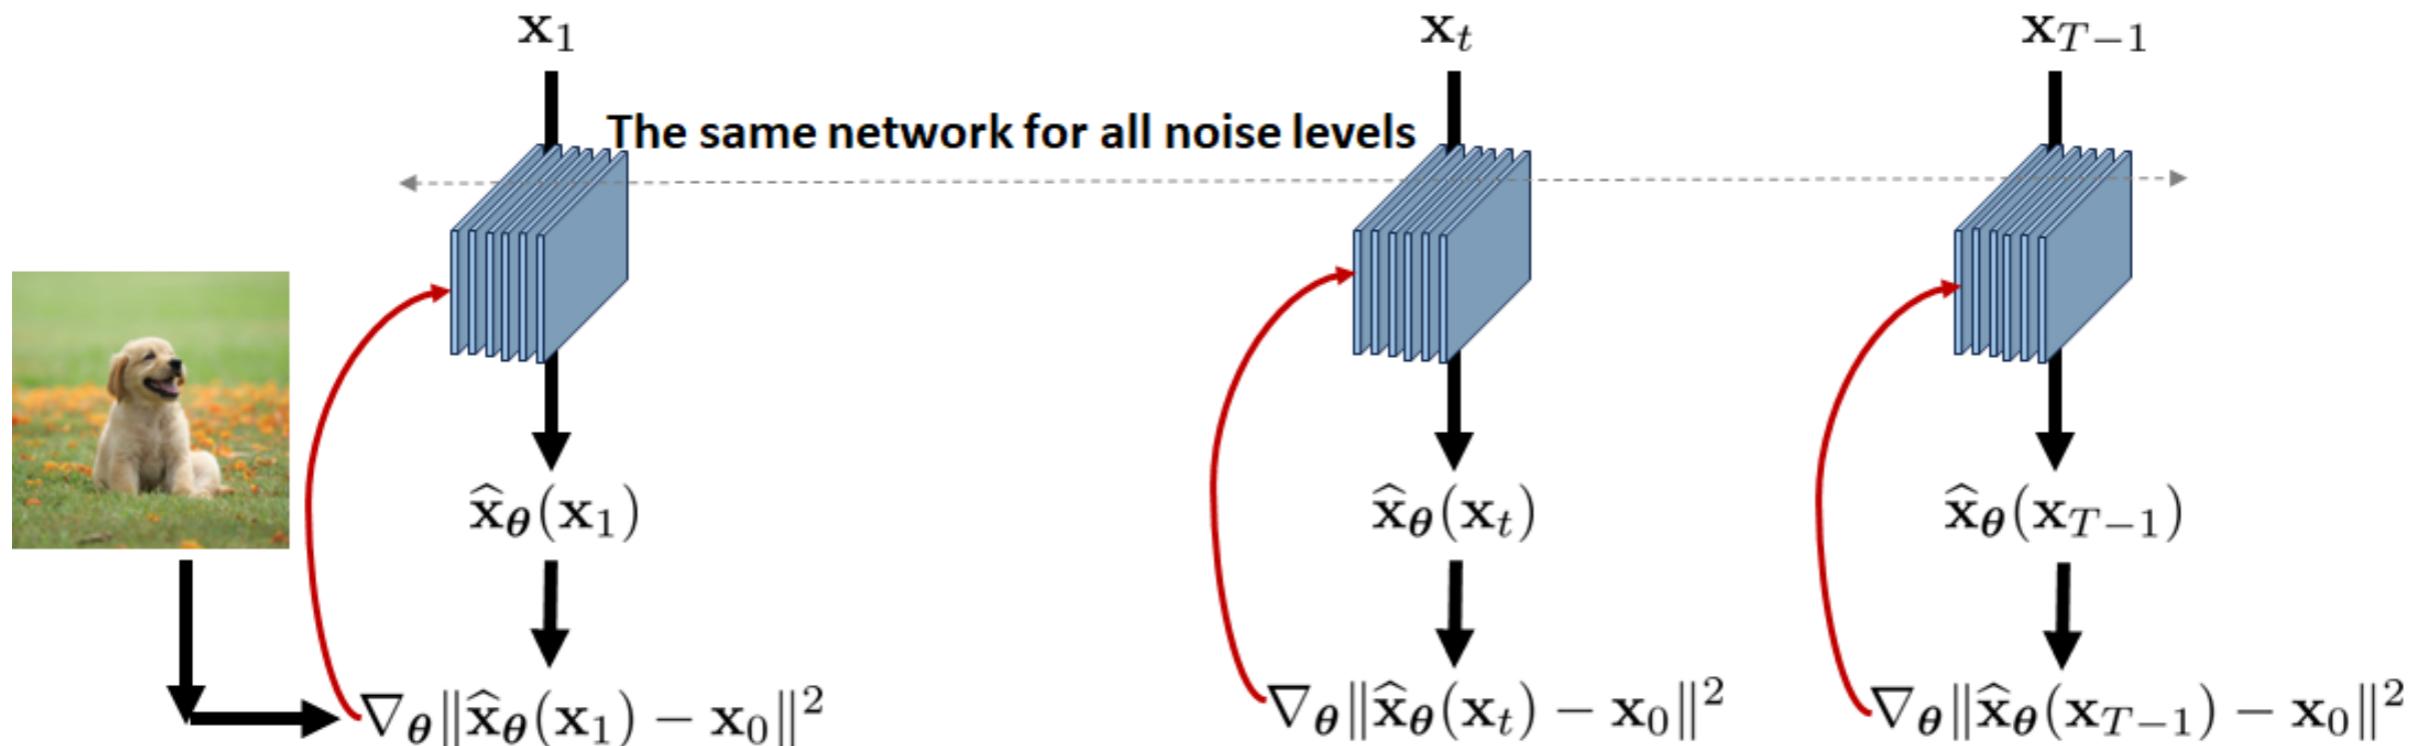
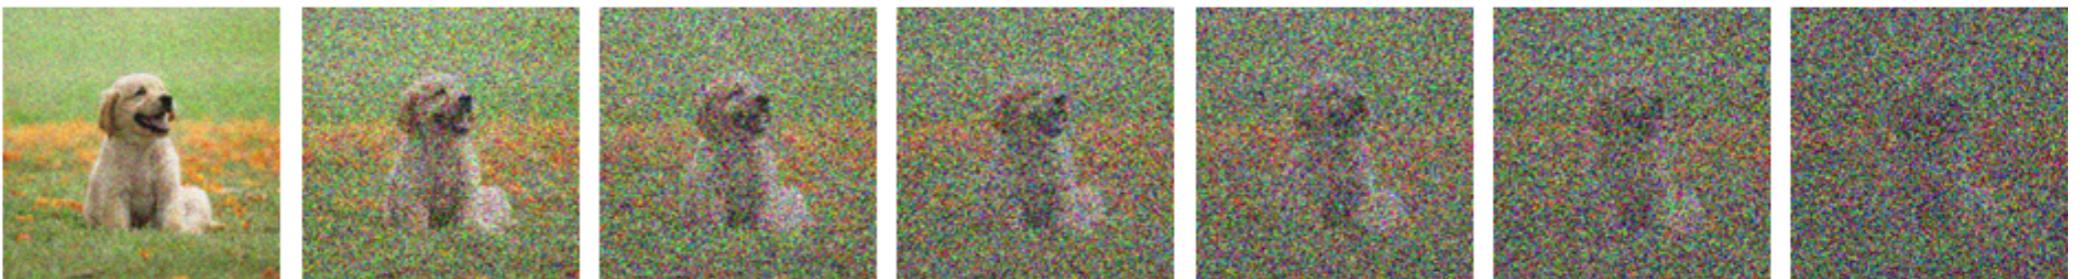
# Diffusion Models

## Training of diffusion models:

After around ~5 pages of algebraic derivations (no conceptual challenges, though):

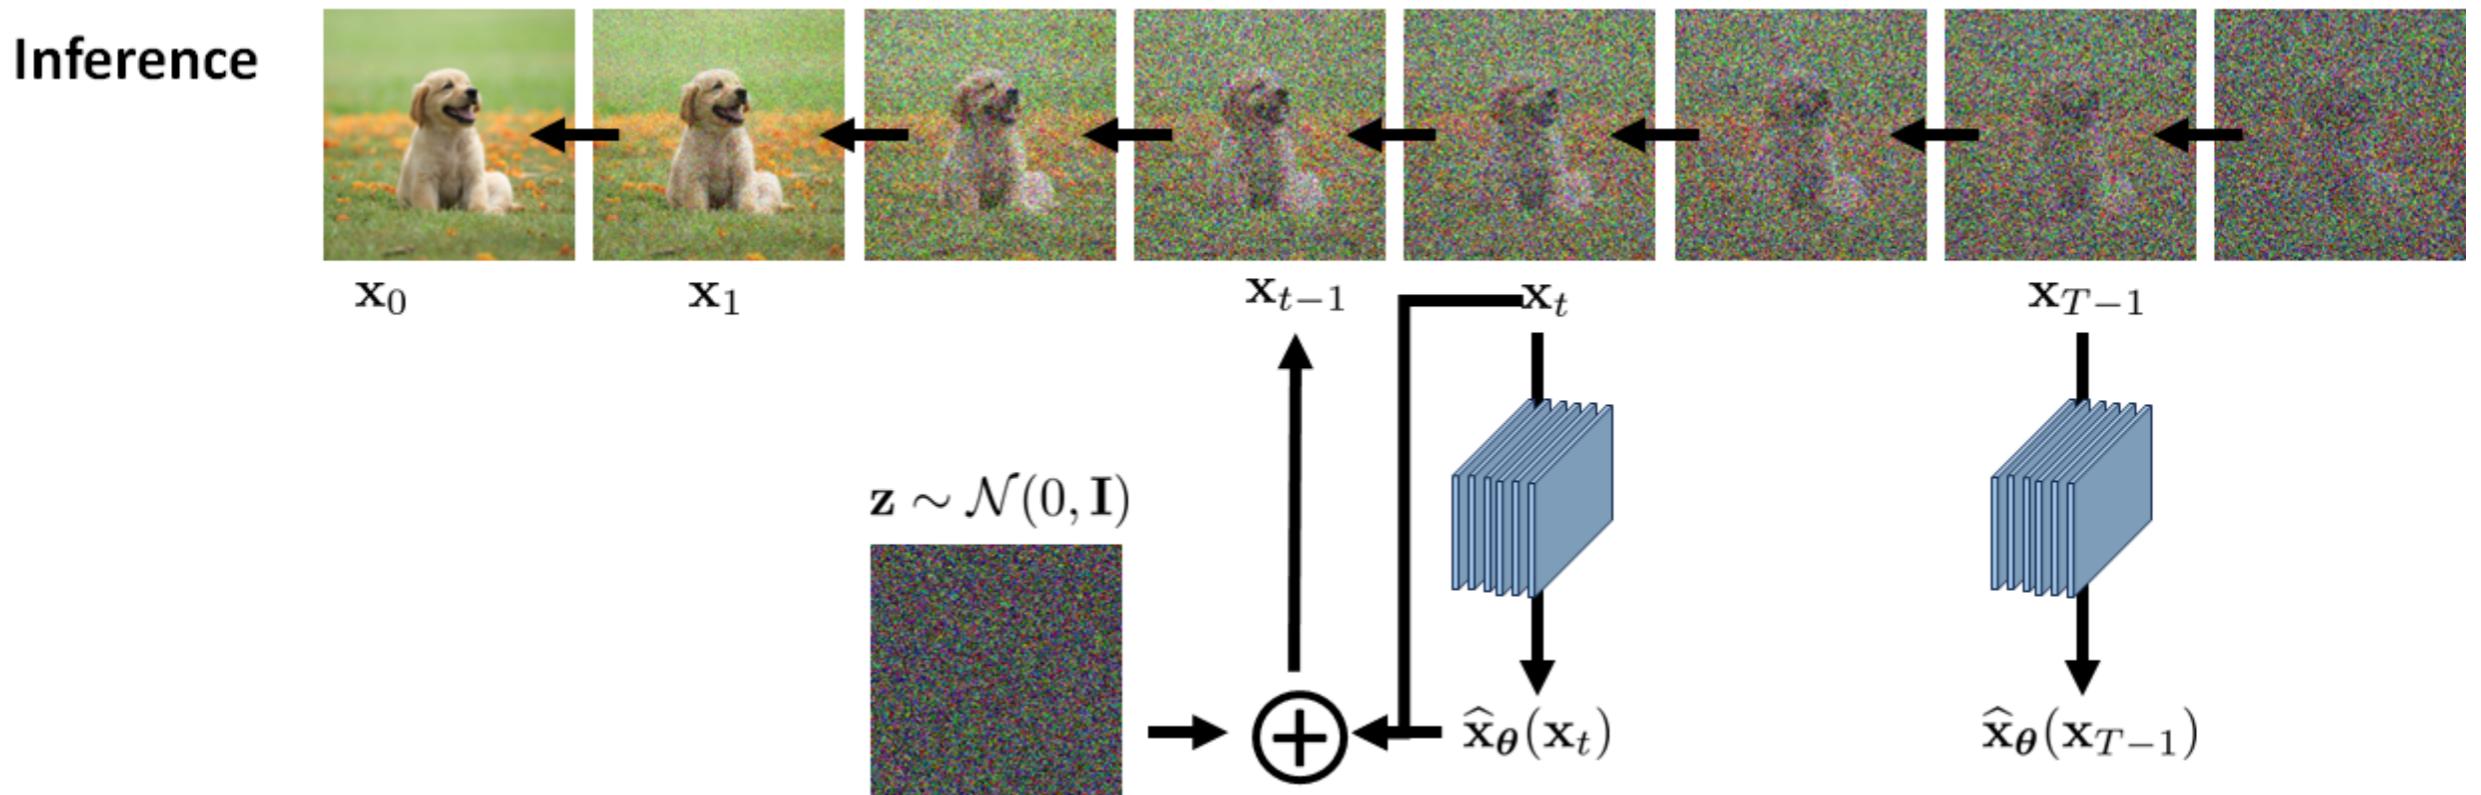
$$\text{ELBO}_\theta(x_0) = - \sum_{t=1}^T \frac{1}{2\sigma^2} \frac{(1 - \alpha_t)\bar{\alpha}_{t-1}}{(1 - \bar{\alpha}_t)^2} \mathbb{E}_{q_\phi(x_t | x_0)} [\|\hat{x}_\theta(x_t) - x_0\|_2^2]$$

Training



# Diffusion Models

## Inference of diffusion models:



- Start from the  $t = T$  step to sample from  $\mathcal{N}(0, I)$
- For  $t = T - 1, \dots, 0$ , gradually apply the trained denoising network to sample from  $p_\theta(x_{t-1} | x_t)$
- Recall that in the backward process we define  $p_\theta(x_{t-1} | x_t)$  to be a Gaussian so the sampling is tractable

# Diffusion Models

---

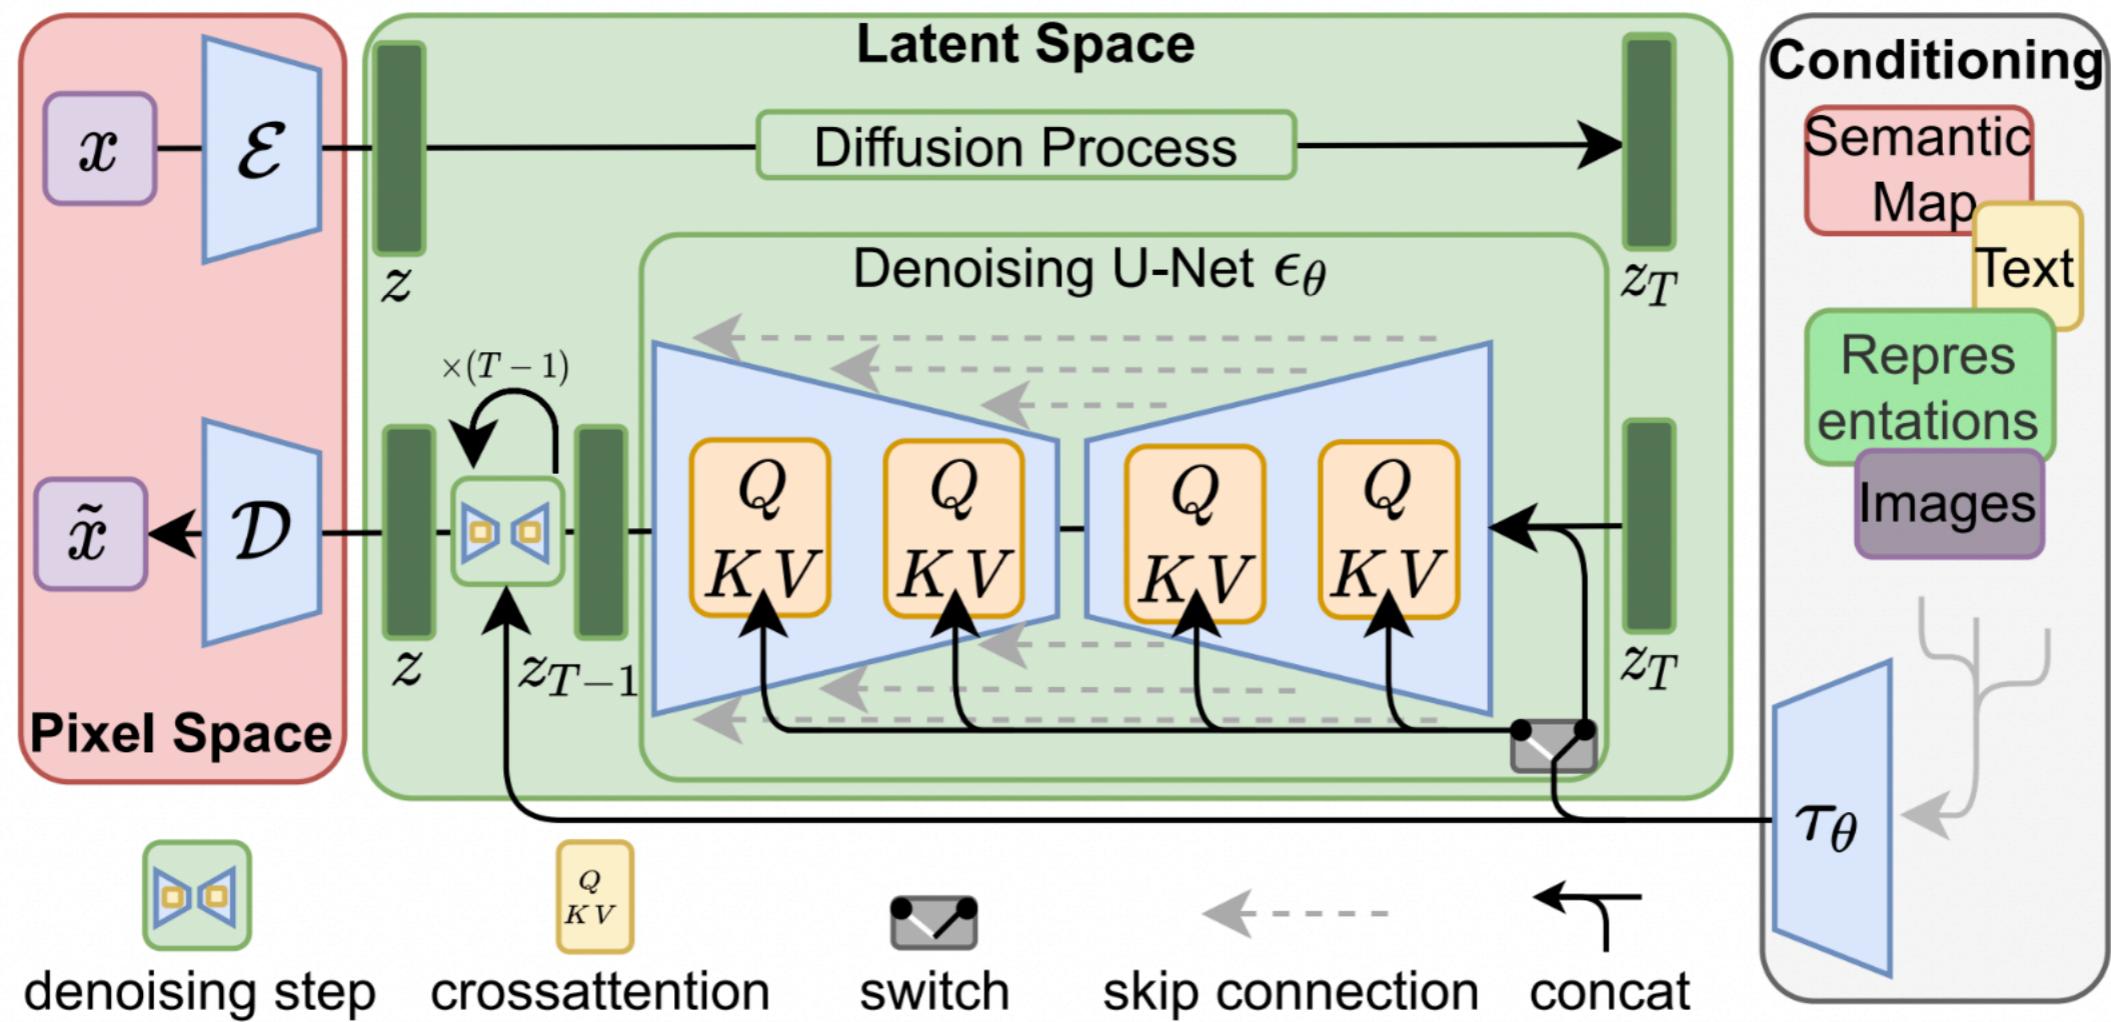
Examples:



# Diffusion Models

## Examples:

Can be combined with textual encoding to enable conditional generation, i.e., text-to-image (Imagen), text-to-video etc.



# Diffusion Models

## Examples:

Can be combined with textual encoding to enable conditional generation, i.e., text-to-image (Imagen). text-to-video etc.



# Diffusion Models

---

## Applications:

- Midjourney
- Huggingface stable diffusion
- Dall-E 3
- .....